

# 3. Linear Neural Networks

Junggwon Shin, [shinjungwon@gmail.com](mailto:shinjungwon@gmail.com)

Summary for Dive Into Deep Learning, [https://d2l.ai/chapter\\_preface/index.html](https://d2l.ai/chapter_preface/index.html)

**3.1 Linear Regression**

**3.2. Linear Regression Implementation from Scratch**

**3.3 Concise Implementation of Linear Regression**

A method for modeling the relationship between one or more **independent variables** and a **dependent variable**

## Assumptions on Linear Regression

- All data are independently and identically distributed
  - All data are randomly sampled from the same distribution independently
  - Data not following IID: time-series data
- Linear relationship between the independent variables  $\mathbf{x}$  and the dependent variable  $y$ 
  - **Inductive bias** (a bias required for generalization) for the model
- Well-behaved **noise** (following a Gaussian distribution)
  - **Not** expect to find a real-world dataset exactly equals the linear relationship due to factors such as measurement error
  - Thus, incorporating a noise term to account for such errors.



## Model

$$\text{price} = w_{\text{area}} \cdot \text{area} + w_{\text{age}} \cdot \text{age} + b.$$


dependent variable  $y$   
(output)
independent variables  $\mathbf{x}$   
(input)

$$\hat{y} = w_1 x_1 + \dots + w_d x_d + b.$$

expressing compactly using a dot product

$$\hat{y} = \mathbf{w}^T \mathbf{x} + b.$$

- $\mathbf{w}$ : weight
- $b$ : bias (also called as offset or intercept)
- Strictly speaking, linear regression is an **affine transformation**
  - a kind of space transformation conserving linearity and parallelism including rotation, reflection, scaling, and others
- **Goal**: Choosing weights  $\mathbf{w}$  and bias  $b$  which optimally fit the given set of pairs between input and output data




Elements to find optimal parameter  $\mathbf{w}$ ,  $b$

- Loss function: a quality measure for some given model
- Stochastic Gradient Descent: a procedure for updating the model to improve its quality

## Loss function

- Quantifying the distance between the real and predicted value of the target
  - Non-negative number where smaller values are better and perfect predictions incur a loss of 0
- Squared error
  - prediction for an data sample  $i$ ,  $\hat{y}^{(i)}$
  - sample target (true label)  $i$ ,  $y^{(i)}$



$$l^{(i)}(\mathbf{w}, b) = \frac{1}{2} (\hat{y}^{(i)} - y^{(i)})^2.$$

Loss for dataset of  $n$  examples

$$L(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n l^{(i)}(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n \frac{1}{2} (\mathbf{w}^T \mathbf{x}^{(i)} + b - y^{(i)})^2.$$

Find optimal parameter that minimize the total loss across all training examples

$$\mathbf{w}^*, b^* = \underset{\mathbf{w}, b}{\operatorname{argmin}} L(\mathbf{w}, b).$$

## Normal Equation

- Analytic solution for the simple liner regression problem

$$\mathbf{y} = X\Theta + \textcircled{\mathbf{e}}^{\text{error}}$$

$$\mathbf{e} = \mathbf{y} - X\Theta$$

Minimize squared error

$$\sum_{j=1}^n \epsilon_j^2 = \mathbf{e}^T \mathbf{e} = (\mathbf{y} - X\Theta)^T (\mathbf{y} - X\Theta)$$

$$= \mathbf{y}^T \mathbf{y} - 2\Theta^T X^T \mathbf{y} + \Theta^T X^T X \Theta$$

Find optimal weight

$$\frac{\partial (\mathbf{e}^T \mathbf{e})}{\partial \Theta} = -2X^T \mathbf{y} + 2X^T X \Theta = \mathbf{0}$$

$$X^T X \Theta = X^T \mathbf{y} \quad \Rightarrow \quad \hat{\Theta} = (X^T X)^{-1} X^T \mathbf{y}$$

### Limitations of Normal Equation

- Hard to calculate a inverse matrix for large-scale data or features,  $O(n^{2.376})$
- No inverse matrix can be existed if collinearity between features exists (singular matrix)


### Gradient Descent

- Iteratively reducing the error by updating the parameters in the direction that incrementally lowers the loss function
- Taking the derivative of the loss function, which is an average of the losses computed on every single example in the dataset

- **Stochastic Gradient Descent**  $\longrightarrow$  for  $i = 1$  to  $n$  :  $\{ \theta_j^{t+1} = \theta_j^t - \alpha(\Theta^T \mathbf{x}_i - y_i)x_i^{(j)} \}$  for every  $j$ 
  - Updating parameters for a gradient of every sampled data
  - Pros: fast calculation speed, small memory requirement
  - Cons: non-stable (largely fluctuating) learning procedure, higher possibility to fall into local minimum

- **Batch Gradient Descent**  $\longrightarrow$   $\theta_j^{t+1} = \theta_j^t - \alpha \sum_{i=1}^n (\Theta^T \mathbf{x}_i - y_i) x_i^{(j)}$  for every  $j$ 
  - update parameters for the mean of calculated gradients for all data
  - Pros: converged into global minimum
  - Cons: low calculation speed, large memory requirement

- **Mini Batch Gradient Descent**
  - update parameters for the mean of calculated gradients for a **subset** of given data



### *Mini Batch Gradient Descent*

- Initializing the values of the model parameters, typically at random
- Randomly sampling a minibatch  $\mathcal{B}$  consisting of a fixed number of training examples
  - $|\mathcal{B}|$  represents the number of examples in each minibatch
- Updating the parameters in the direction of the negative gradient
  - $\partial_i$  denotes the partial derivative of parameter  $I$
  - $\eta$  denotes the learning rate

$$\begin{aligned}\mathbf{w} &\leftarrow \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{\mathbf{w}} l^{(i)}(\mathbf{w}, b) = \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)}) , \\ b &\leftarrow b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_b l^{(i)}(\mathbf{w}, b) = b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)}) .\end{aligned}$$

### *Hyperparameter Tuning*

- These parameters that are **tunable** but **not updated in the training loop**
- Typically adjust hyperparameters based on the results of the training loop as assessed on a separate **validation dataset**
- $|\mathcal{B}|$  and  $\eta$  for the simple linear regression model

### *Stopping Training*

- Finishing training after certain number of iterations or until some stopping criteria met
- The trained parameters **will not exact minimizers** of the loss because it cannot achieve it exactly in a finite number of steps


### Mini Batch Gradient Descent

- Initializing the values of the model parameters, typically at random
- Randomly sampling a minibatch  $\mathcal{B}$  consisting of a fixed number of training examples
  - $|\mathcal{B}|$  represents the number of examples in each minibatch
- Updating the parameters in the direction of the negative gradient
  - $\partial_i$  denotes the partial derivative of parameter  $I$
  - $\eta$  denotes the learning rate

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{\mathbf{w}} l^{(i)}(\mathbf{w}, b) = \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)}),$$


$$b \leftarrow b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_b l^{(i)}(\mathbf{w}, b) = b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)}).$$

Loss surface of a simple linear regression model



### Loss function of Linear Regression and Deep Neural Networks

- Linear regression works for a learning problem where there is only one minimum over the entire domain
  - The bowl-shaped loss-function
  - Model with simple inductive bias  $\rightarrow$  model with **high bias error** for complex data
- DNN model contains loss surfaces with many minima
  - The loss function of DNN is not bowl-shaped and not convex (much more complex)
  - Model with complex inductive bias  $\rightarrow$  model with **low bias error** for complex data



Loss surface of a DNN model

## Bias-Variance Trade-off

- Kinds of prediction error
  - Bias error: generated due to erroneous hypothesis on a model
    - high bias = underfitting
  - Variance error: from sensitivity to small fluctuations on different input data
    - high variance = overfitting
  - Irreducible error: noise in data

## Bias-variance decomposition of mean squared error

Target output      Prediction result

$$\text{MSE} = E[(y - \hat{f})^2]$$

zero-mean noise in the target output  $\sim N(0, \sigma)$

$$= E[(f + \epsilon - \hat{f})^2]$$

$$= E[(f + \epsilon - \hat{f} + E[\hat{f}] - E[\hat{f}])^2]$$

$$= E[(f - E[\hat{f}])^2] + E[\epsilon^2] + E[(E[\hat{f}] - \hat{f})^2] + 2E[(f - E[\hat{f}])\epsilon] + 2E[\epsilon(E[\hat{f}] - \hat{f})] + 2E[(E[\hat{f}] - \hat{f})(f - E[\hat{f}])]$$

$$= E[(f - E[\hat{f}])^2] + E[\epsilon^2] + E[(E[\hat{f}] - \hat{f})^2] + 2E[\epsilon]E[(f - E[\hat{f}])] + 2E[\epsilon]E[E[\hat{f}] - \hat{f}] + 2E[(E[\hat{f}] - \hat{f})(f - E[\hat{f}])]$$

$$= E[(f - E[\hat{f}])^2] + E[\epsilon^2] + E[(E[\hat{f}] - \hat{f})^2] + 0 + 0 + 0$$

zero-mean

$E[XY] = E[X]E[Y]$   
Mean of deviation = 0



**Bias<sup>2</sup>**   **Irreducible**   **Variance**

Fluctuation (variance) of the prediction results

Error between true result  
and expected result from model

**Error**

Variance of the random noise,  $E[(\epsilon - 0)^2] = \sigma^2$



## Statistical Interpretation of Linear Regression

- *Dose least-squares loss function find actual optimal parameters of a linear regression model?*


MLE for weight parameters of a linear regression model with gaussian noise equals Minimizing MSE of the model

$$\begin{aligned}
 y_i &= \Theta \mathbf{x}_i + \epsilon_i \quad \text{Gaussian noise} \\
 p(\epsilon_i) &= \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\epsilon_i^2}{2\sigma^2}\right) \\
 p(y_i|\mathbf{x}_i; \theta) &= \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \Theta \mathbf{x}_i)^2}{2\sigma^2}\right) \\
 \text{MLE for weight parameter } \Theta & \rightarrow L(\Theta) = L(\Theta; X, Y) = p(Y|X; \Theta) \\
 L(\Theta) &= \prod_{i=1}^n p(y_i|\mathbf{x}_i; \Theta) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \Theta \mathbf{x}_i)^2}{2\sigma^2}\right) \quad \text{IID assumptions} \\
 \text{Log likelihood} & \rightarrow \log L(\Theta) = \log \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \Theta \mathbf{x}_i)^2}{2\sigma^2}\right) = \sum_{i=1}^n \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \Theta \mathbf{x}_i)^2}{2\sigma^2}\right) \\
 \text{Maximizing log likelihood} & = n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \frac{1}{2} \sum_{i=1}^n (y_i - \Theta^T \mathbf{x}_i)^2 \quad \text{Minimizing Squared error} \\
 & \quad \sigma \text{ is some fixed constant}
 \end{aligned}$$

## Interpretation of LR model in a perspective on Deep Neural Network


Linear regression model is A kind of fully-connected layer or dense layer

- Every input is connected to every output



Layer representation of a liner regression model

$$o_1 = w_1 x_1 + w_2 x_2 + \dots + w_d x_d$$



axon

nucleus

dendrites

$$y = \sum_i x_i w_i + b,$$



### Further material for following chapters

[https://github.com/howawindelu/dive-into-deep-learning/blob/master/week3/week3\\_1\\_implementation\\_lukeshin.ipynb](https://github.com/howawindelu/dive-into-deep-learning/blob/master/week3/week3_1_implementation_lukeshin.ipynb)

- 3.1.2. Vectorization for Speed
- 3.2. Linear Regression Implementation from Scratch
- 3.3 Concise Implementation of Linear Regression