

Dive into Deep Learning

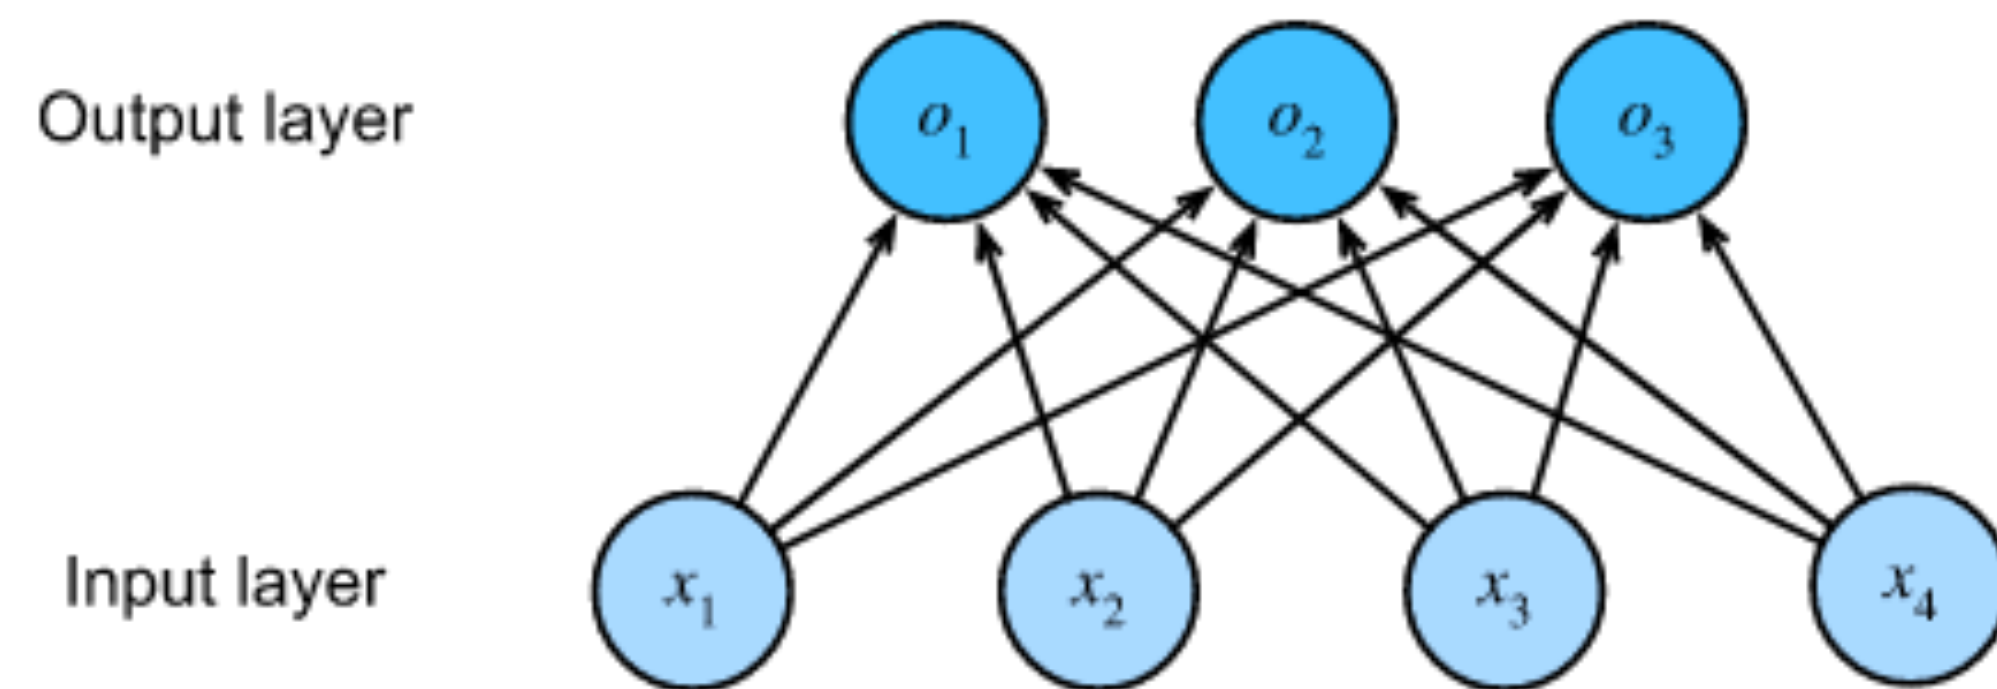
Chapter 4. Multilayer Perceptrons

Wayne L, Oct 18, 2020

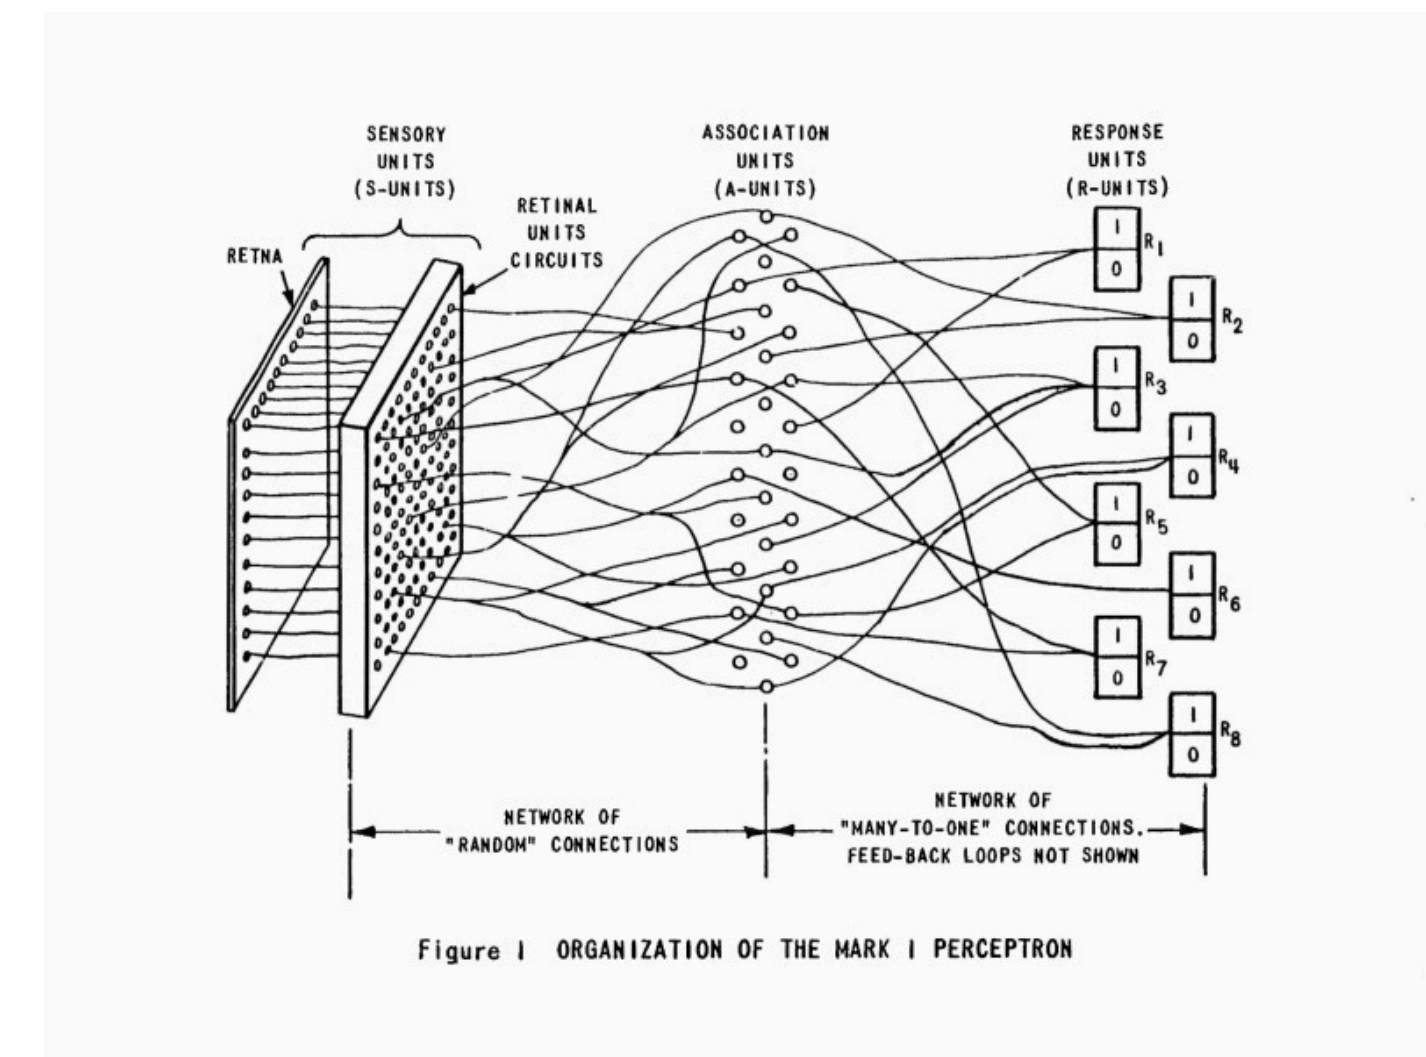
Previously

Perceptron (Linear = Single = Simple = Feedforward)

- This model mapped our inputs directly to our outputs via single affine transformation, followed by a softmax operation.



$$y = \mathbf{w}\mathbf{x} + b$$




Mark 1 Perceptron (1959), Frank Rosenblatt

Previously

AI History


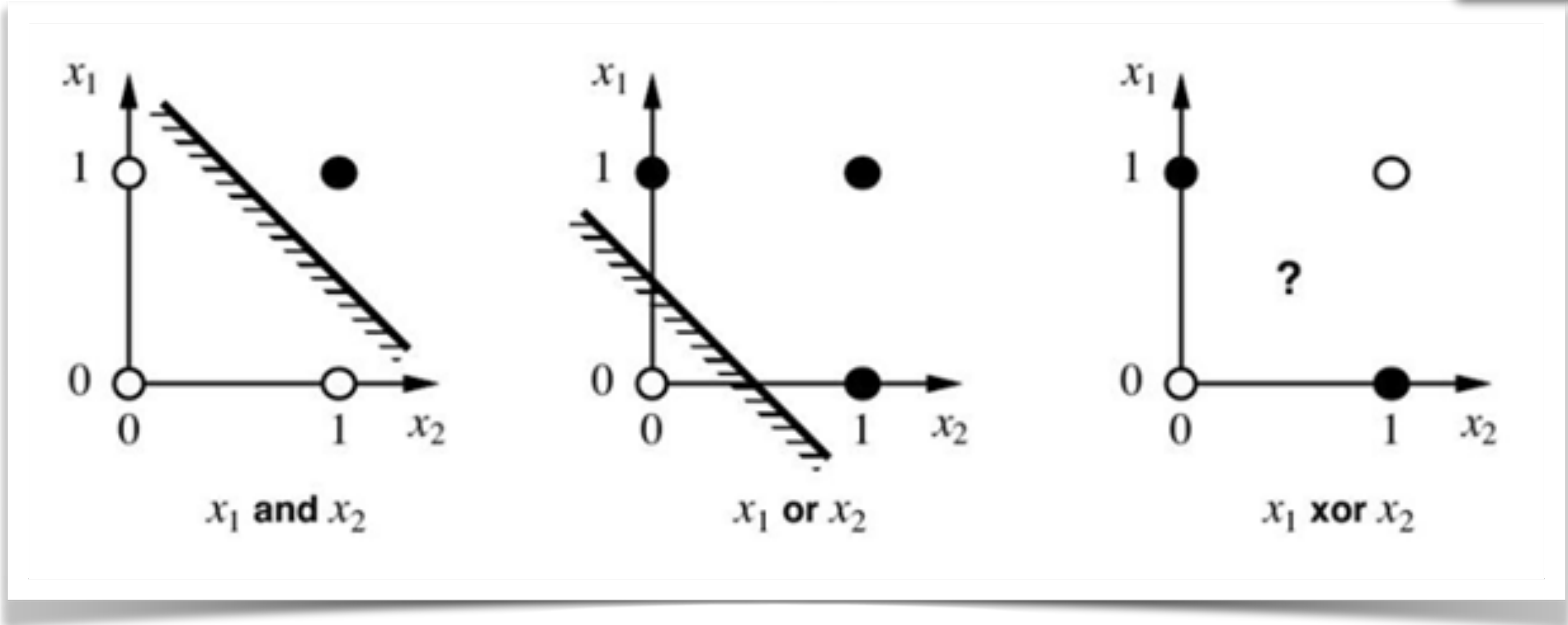
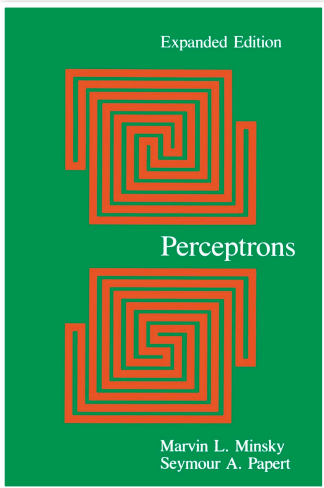
AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING” ...



Timeline of AI Development
▪ 1950s-1960s: First AI boom - the age of reasoning, prototype AI developed
▪ 1970s: AI winter I
▪ 1980s-1990s: Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
▪ 1990s: AI winter II
▪ 1997: Deep Blue beats Gary Kasparov
▪ 2006: University of Toronto develops Deep Learning
▪ 2011: IBM's Watson won Jeopardy
▪ 2016: Go software based on Deep Learning beats world's champions

XOR

Perceptrons (1969), Marvin Minsky




No one on earth had found a viable way to train

Multilayer Perceptrons

Perceptrons (Multi = Deep feedforward)

- The goal of a feedforward network is to approximate some function f^* .
- For example, for a classifier, $y = f^*(x)$ maps an input x to a category y . A feedforward network defines a mapping $\mathbf{y} = f(\mathbf{x}; \theta)$ and learns the value of the parameters θ that result in the best function approximation.
- **Universal approximation theorem** means that regardless of what function we are trying to learn, we know that a large MLP will be able to represent this function.

$$f(x) = f^*(x) \approx y$$



$$\mathbf{O} = \mathbf{H}\mathbf{W}^{(2)} + \mathbf{b}^{(2)}$$

$$\mathbf{H} = \mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)}$$

$$\begin{aligned}\mathbf{W} &= \mathbf{W}^{(1)}\mathbf{W}^{(2)} \\ \mathbf{b} &= \mathbf{b}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)}\end{aligned}$$

$$\mathbf{O} = (\mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)})\mathbf{W}^{(2)} + \mathbf{b}^{(2)} = \mathbf{X}\mathbf{W}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)} = \mathbf{X}\mathbf{W} + \mathbf{b}$$

Even though MLP is going deeper, it can be equivalent with single-layer model !

Multilayer Perceptrons

From Linear to Nonlinear

- In order to realize the potential of multilayer architectures, we need one more key ingredient: a nonlinear *activation function* σ to be more **expressive**.
- MLPs are **universal approximators**, however, it does not mean that we can solve all of problems with MLPs. In fact, we can approximate many functions much more compactly by using deeper (or wider) networks.
- Each neuron acts as a **linear SVM**, however, ...
 - its output is not interpreted immediately,
 - but it becomes a new feature,
 - to be forwarded to the next layer for further analysis **#SVMcascade**

$$\mathbf{H}^{(1)} = \sigma_1(\mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)})$$

$$\mathbf{H}^{(2)} = \sigma_2(\mathbf{H}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)})$$

Multilayer Perceptrons

Activation function

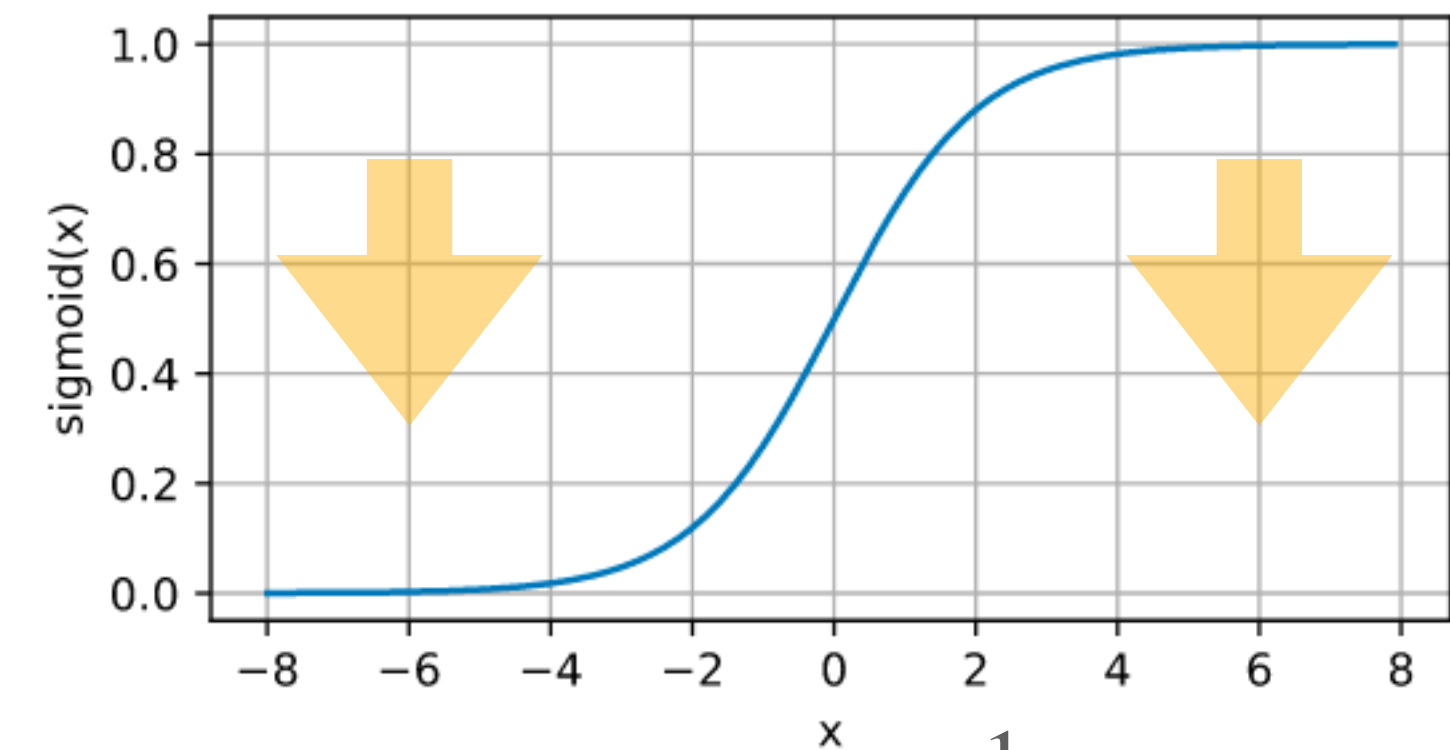
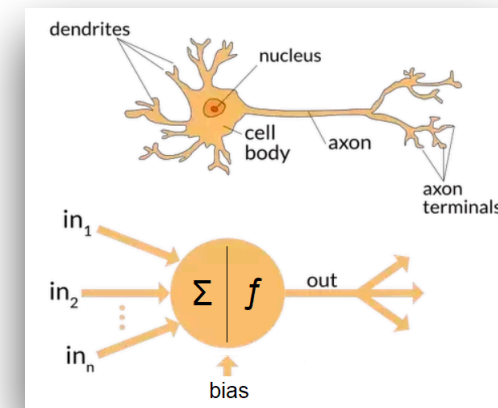
- **Activation function** decide whether a neuron should be activated or not by calculating the weighted sum and further adding bias with it.
- They are **differentiable** operators to transform input signal to outputs, while most of them add non-linearity.

<code>nn.ELU</code>	Applies the element-wise function:
<code>nn.Hardshrink</code>	Applies the hard shrinkage function element-wise:
<code>nn.Hardsigmoid</code>	Applies the element-wise function:
<code>nn.Hardtanh</code>	Applies the HardTanh function element-wise
<code>nn.Hardswish</code>	Applies the hardswish function, element-wise, as described in the paper:
<code>nn.LeakyReLU</code>	Applies the element-wise function:
<code>nn.LogSigmoid</code>	Applies the element-wise function:
<code>nn.MultiheadAttention</code>	Allows the model to jointly attend to information from different representation subspaces.
<code>nn.PReLU</code>	Applies the element-wise function:
<code>nn.ReLU</code>	Applies the rectified linear unit function element-wise:
<code>nn.ReLU6</code>	Applies the element-wise function:
<code>nn.RReLU</code>	Applies the randomized leaky rectified liner unit function, element-wise, as described in the paper:
<code>nn.SELU</code>	Applied element-wise, as:
<code>nn.CELU</code>	Applies the element-wise function:
<code>nn.GELU</code>	Applies the Gaussian Error Linear Units function:
<code>nn.Sigmoid</code>	Applies the element-wise function:
<code>nn.Softplus</code>	Applies the element-wise function:
<code>nn.Softshrink</code>	Applies the soft shrinkage function elementwise:
<code>nn.Softsign</code>	Applies the element-wise function:
<code>nn.Tanh</code>	Applies the element-wise function:
<code>nn.Tanhshrink</code>	Applies the element-wise function:
<code>nn.Threshold</code>	Thresholds each element of the input Tensor.

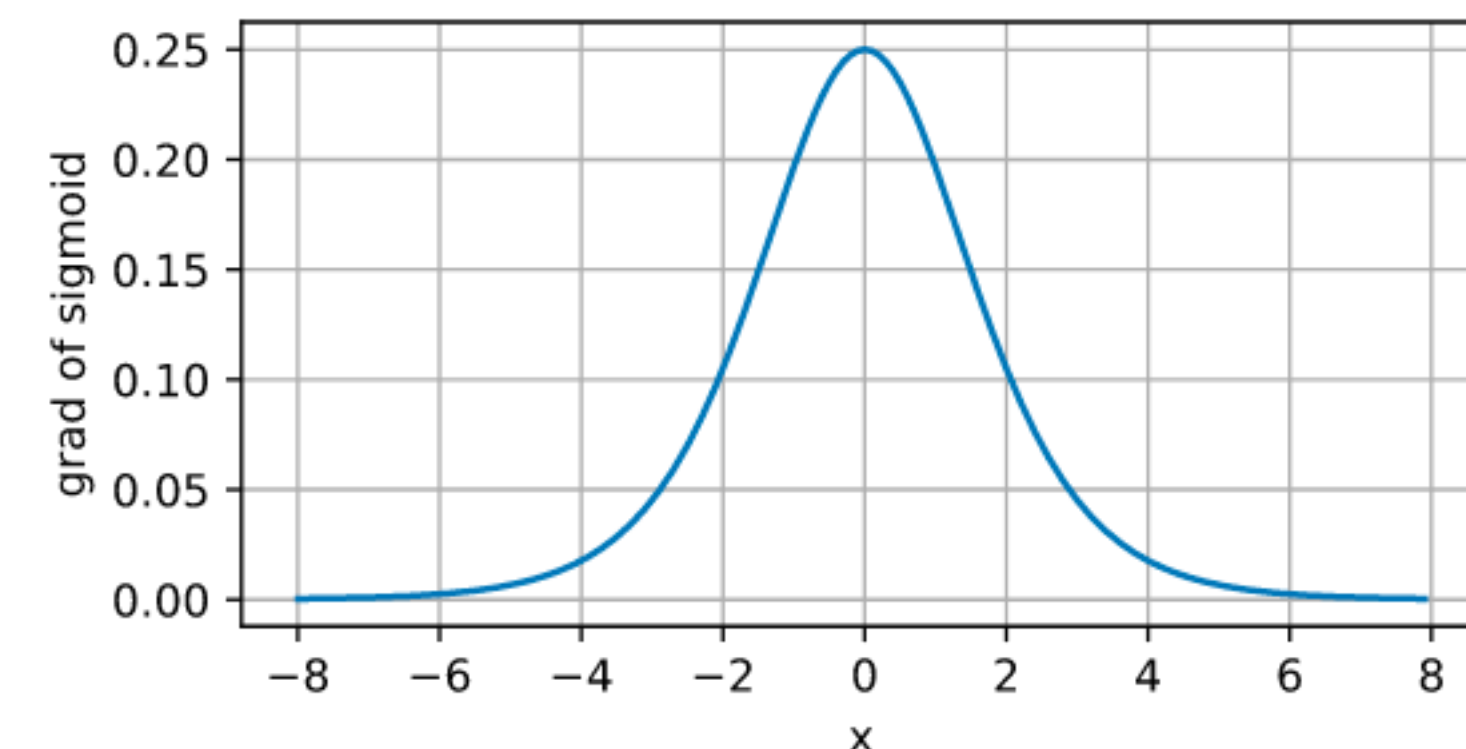
Multilayer Perceptrons

Sigmoid

- In the earliest neural networks, scientists were interested in modeling biological neurons which either fire or do not fire.
- When attention shifted to gradient based learning, the sigmoid function was a natural choice because it is a smooth, differentiable approximation to a thresholding unit.
- widely used as activation functions on the output units, when we want to interpret the outputs as probabilities for binary classification problems (special case of the softmax).



$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}$$



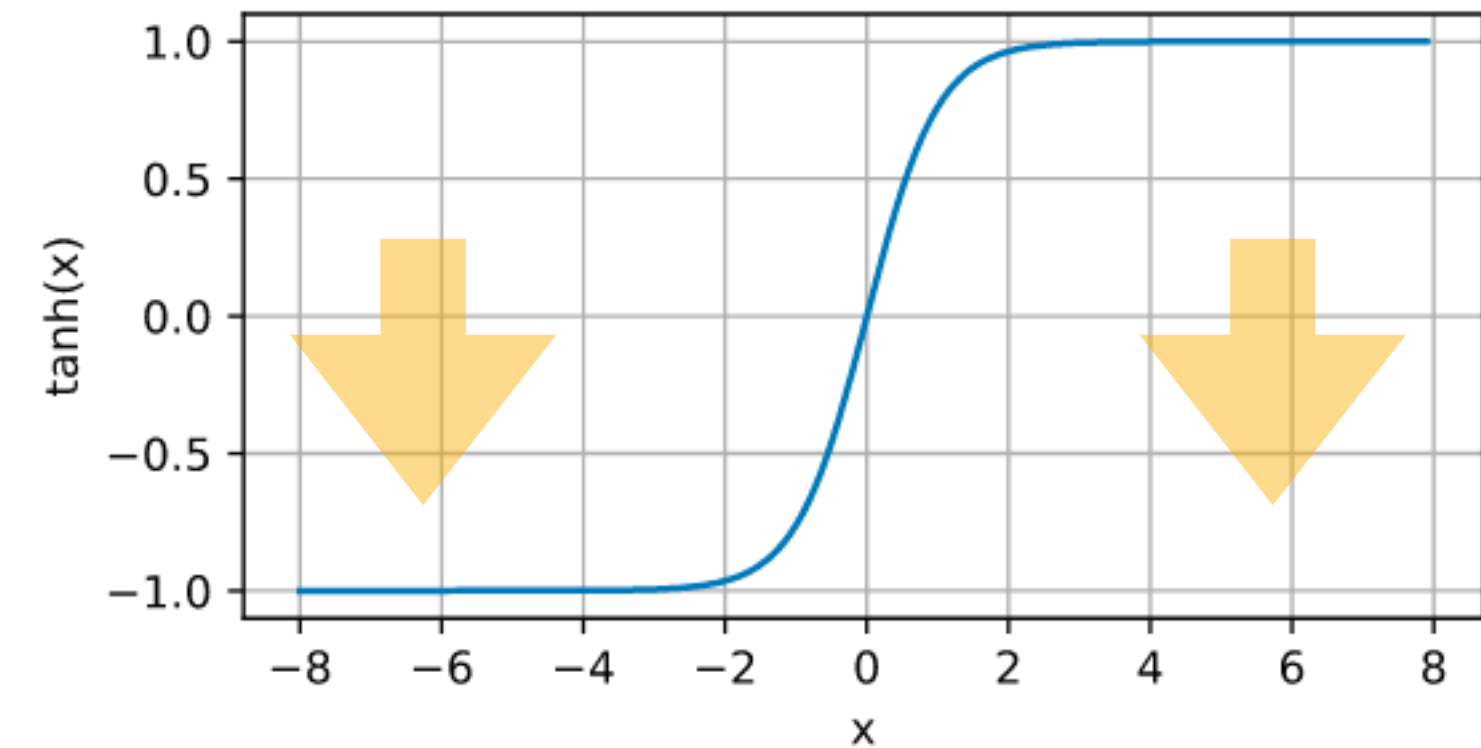
vanishing
gradient

$$\frac{d}{dx} \text{sigmoid}(x) = \frac{\exp(-x)}{(1 + \exp(-x))^2} = \text{sigmoid}(x)(1 - \text{sigmoid}(x))$$

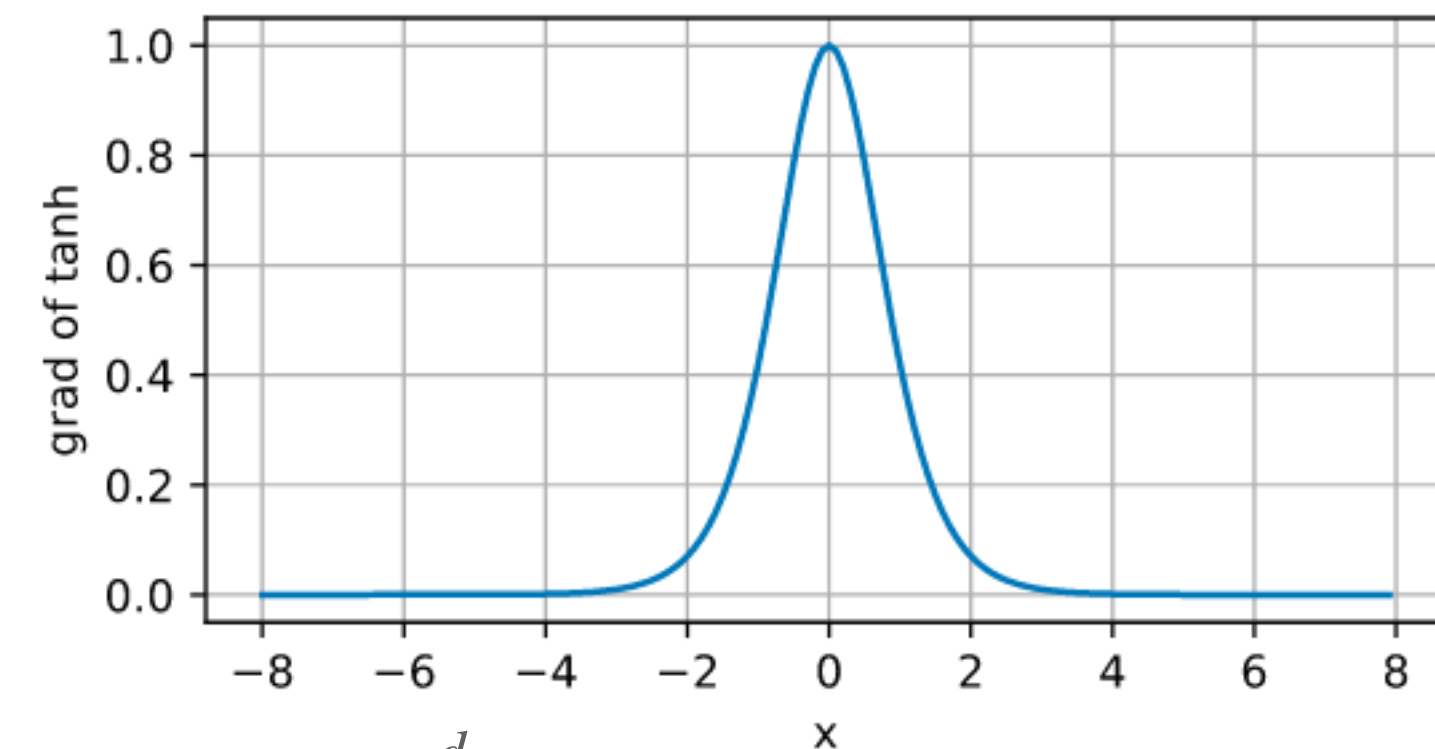
Multilayer Perceptrons

Tanh (hyperbolic tangent)

- Like the sigmoid function, the tanh function also squashes its inputs, transforming them into elements on the interval between -1 and 1.
- Although the shape of the function is similar to that of the sigmoid function, the tanh function exhibits point symmetry about the origin of the coordinate system.



$$\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$$



$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x)$$

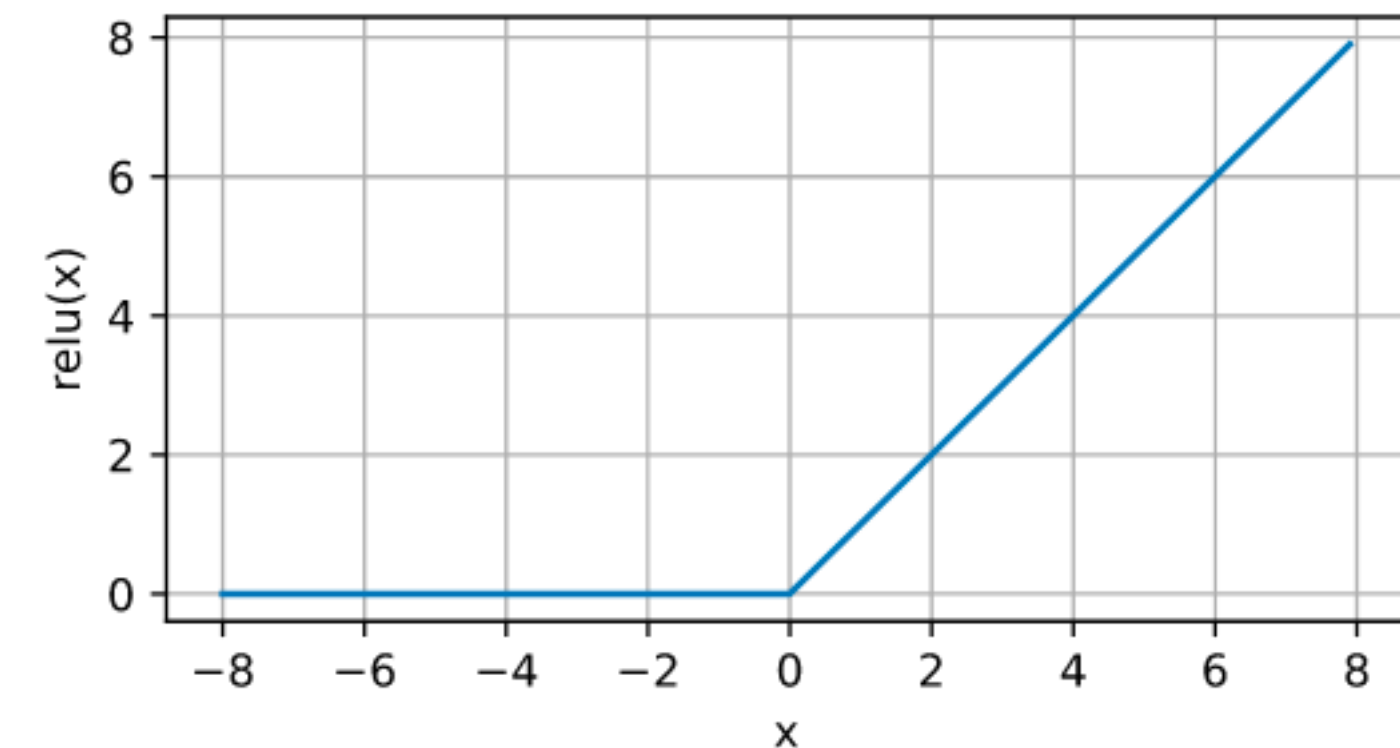
vanishing
gradient

Multilayer Perceptrons

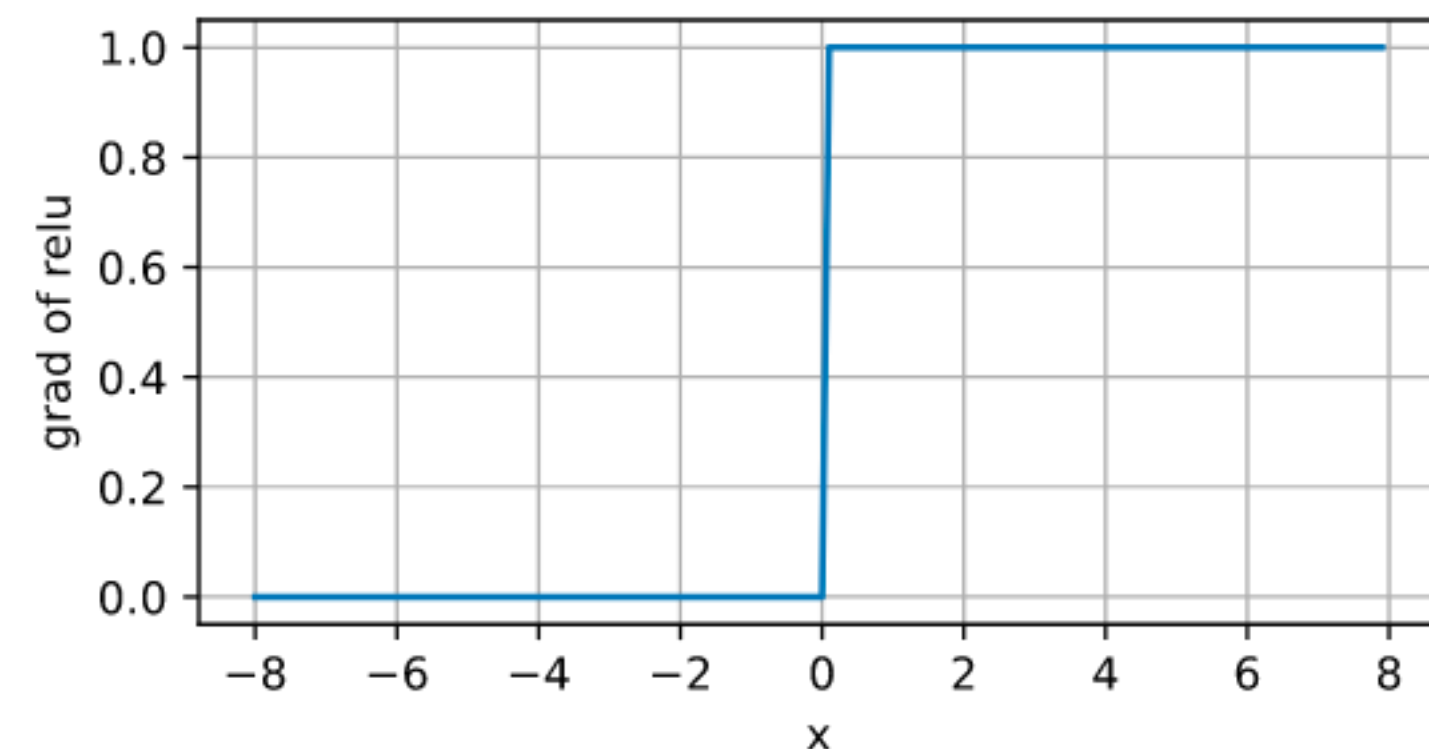
ReLU (rectified linear unit)

- The reason for using ReLU is that its derivatives are particularly well behaved: either they vanish or they just let the argument through.
- This makes optimization better behaved and it mitigated the well-documented problem of vanishing gradients that plagued previous versions of neural networks.

$$pReLU(x) = \max(0, x) + \alpha \min(0, x)$$



$$ReLU(x) = \max(x, 0)$$

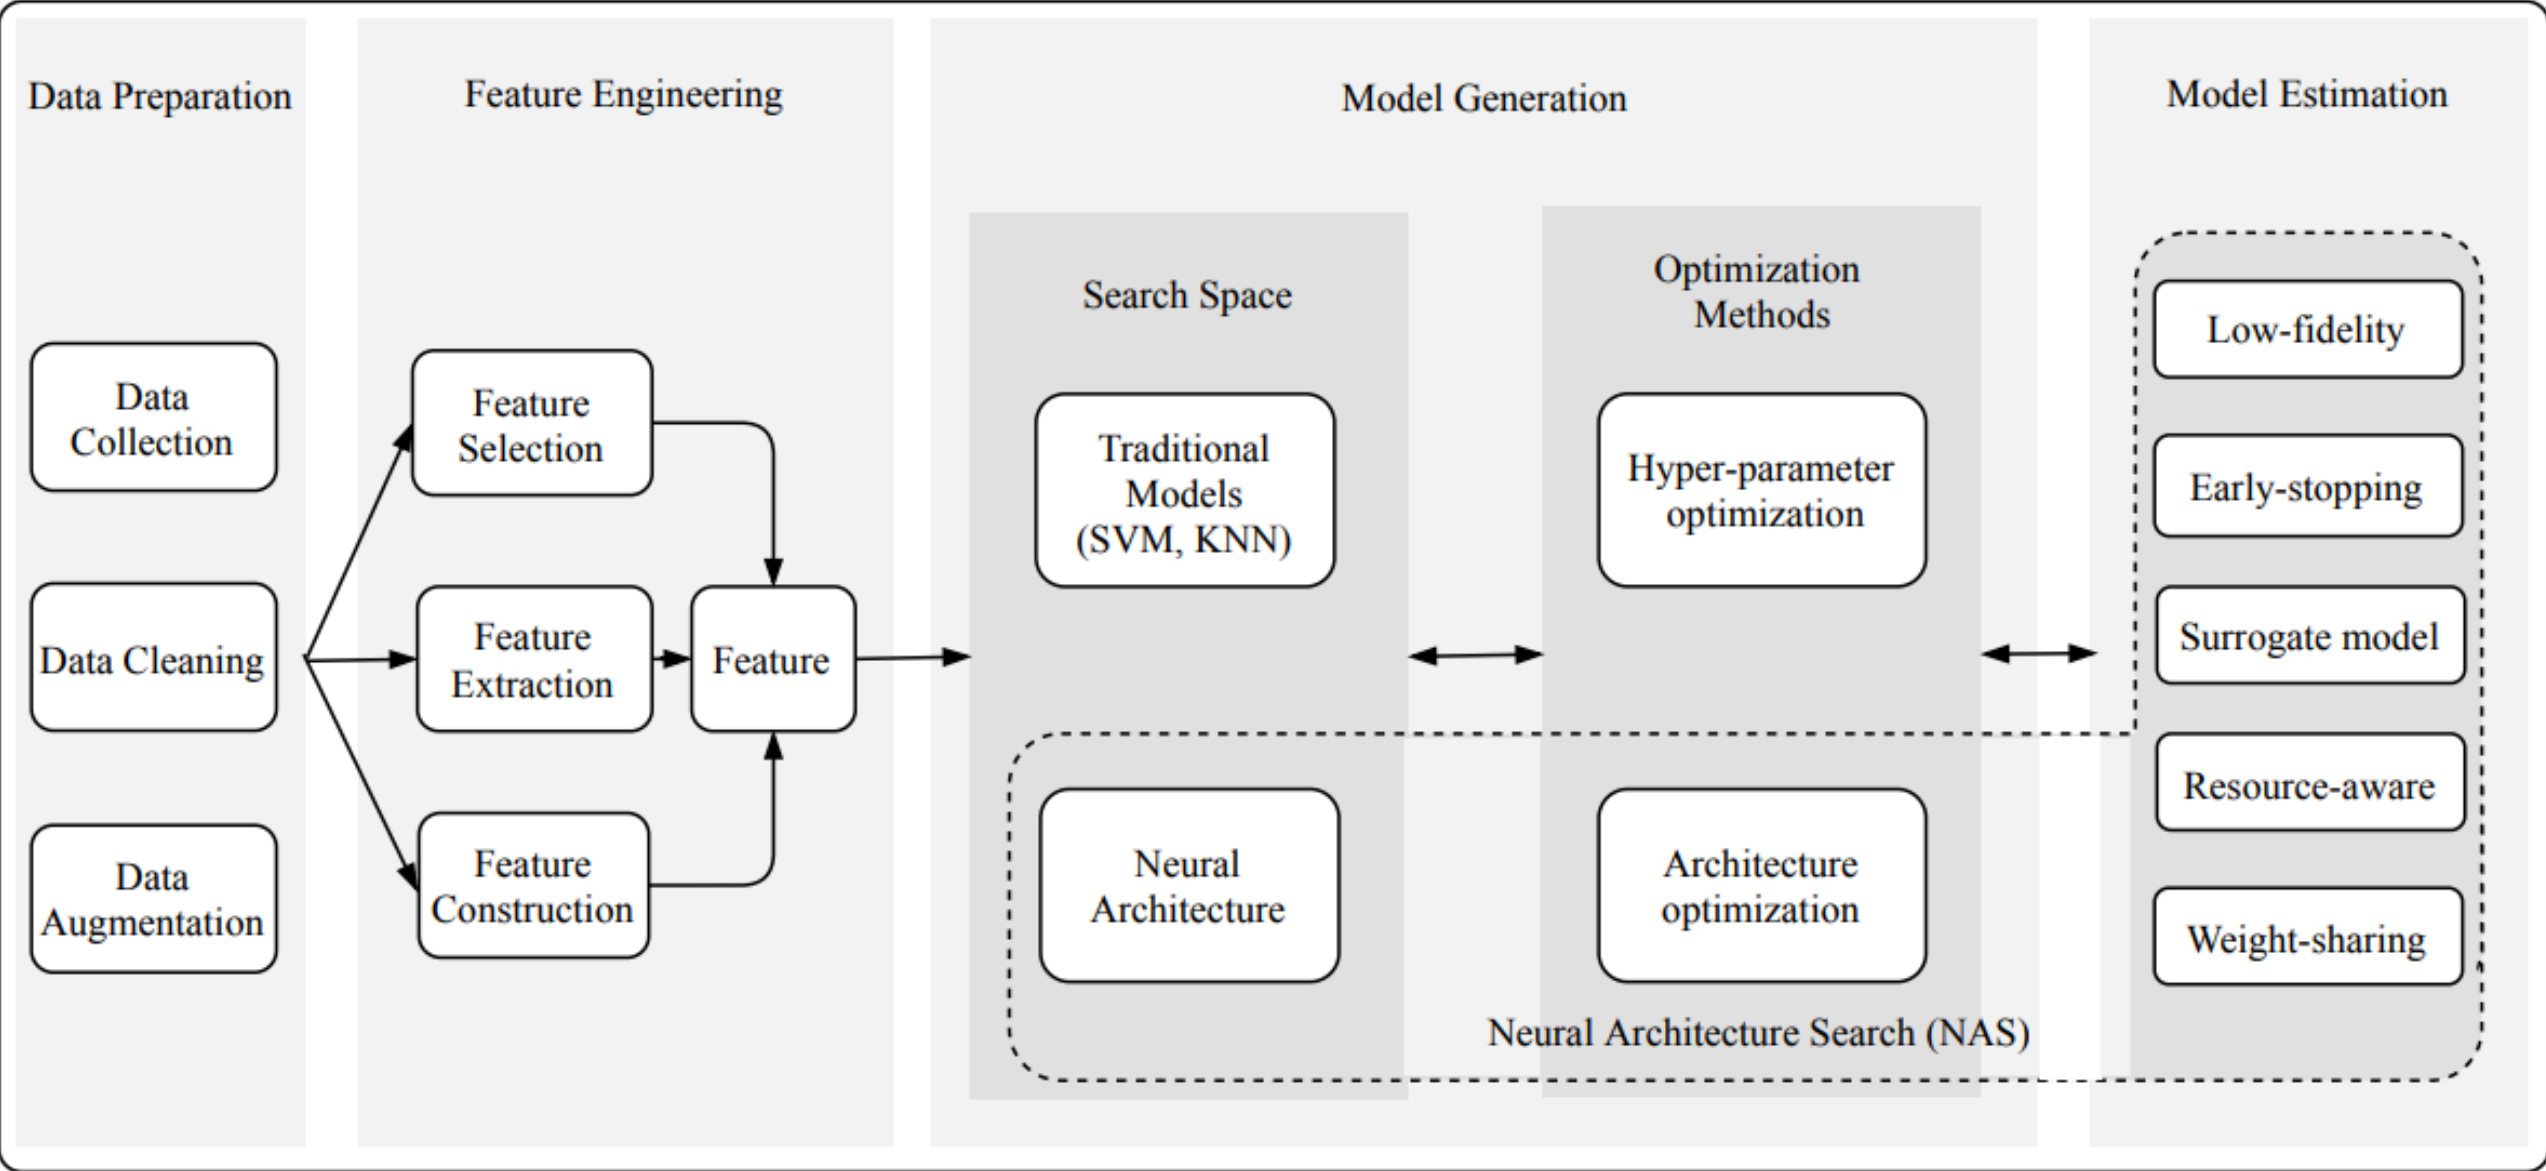


“Predicting in advance which will work best is usually impossible.”

Deep Learning (Chapter 6) , Ian Goodfellow et al.

Auto ML

- #AutoML: automated machine learning
- #NAS: neural architecture search
- #HPO: hyper-parameter optimization



- (1) <https://github.com/microsoft/nni>
- (2) <https://github.com/automl/auto-sklearn>
- (3) <https://github.com/google/automl>

CIFAR 10 (<https://www.cs.toronto.edu/~kriz/cifar.html>)

Reference	Published in	#Params (Millions)	Top-1 Acc(%)	GPU Days	#GPUs	AO
ResNet-110 [2]	ECCV16	1.7	93.57	-	-	Maunally designed
PyramidNet [205]	CVPR17	26	96.69	-	-	
DenseNet [124]	CVPR17	25.6	96.54	-	-	
GeNet#2 (G-50) [30]	ICCV17	-	92.9	17	-	EA
Large-scale ensemble [25]	ICML17	40.4	95.6	2,500	250	
Hierarchical-EAS [19]	ICLR18	15.7	96.25	300	200	
CGP-ResSet [28]	IJCAI18	6.4	94.02	27.4	2	
AmoebaNet-B (N=6, F=128)+c/o [26]	AAAI19	34.9	97.87	3,150	450 K40	
AmoebaNet-B (N=6, F=36)+c/o [26]	AAAI19	2.8	97.45	3,150	450 K40	
Lemonade [27]	ICLR19	3.4	97.6	56	8 Titan	
EENA [146]	ICCV19	8.47	97.44	0.65	1 Titan Xp	
EENA (more channels)[146]	ICCV19	54.14	97.79	0.65	1 Titan Xp	
NASv3[12]	ICLR17	7.1	95.53	22,400	800 K40	RL
NASv3+more filters [12]	ICLR17	37.4	96.35	22,400	800 K40	
MetaQNN [23]	ICLR17	-	93.08	100	10	
NASNet-A (7 @ 2304)+c/o [15]	CVPR18	87.6	97.60	2,000	500 P100	
NASNet-A (6 @ 768)+c/o [15]	CVPR18	3.3	97.35	2,000	500 P100	
Block-QNN-Connection more filter [16]	CVPR18	33.3	97.65	96	32 1080Ti	
Block-QNN-Depthwise, N=3 [16]	CVPR18	3.3	97.42	96	32 1080Ti	
ENAS+macro [13]	ICML18	38.0	96.13	0.32	1	
ENAS+micro+c/o [13]	ICML18	4.6	97.11	0.45	1	
Path-level EAS [136]	ICML18	5.7	97.01	200	-	
Path-level EAS+c/o [136]	ICML18	5.7	97.51	200	-	
ProxylessNAS-RL+c/o[129]	ICLR19	5.8	97.70	-	-	
FPNAS[206]	ICCV19	5.76	96.99	-	-	
DARTS(first order)+c/o[17]	ICLR19	3.3	97.00	1.5	4 1080Ti	GD
DARTS(second order)+c/o[17]	ICLR19	3.3	97.23	4	4 1080Ti	
sharpDARTS [175]	ArXiv19	3.6	98.07	0.8	1 2080Ti	
P-DARTS+c/o[125]	ICCV19	3.4	97.50	0.3	-	
P-DARTS(large)+c/o[125]	ICCV19	10.5	97.75	0.3	-	
SETN[207]	ICCV19	4.6	97.31	1.8	-	
GDAS+c/o [151]	CVPR19	2.5	97.18	0.17	1	
SNAS+moderate constraint+c/o [152]	ICLR19	2.8	97.15	1.5	1	
BayesNAS[208]	ICML19	3.4	97.59	0.1	1	
ProxylessNAS-GD+c/o[129]	ICLR19	5.7	97.92	-	-	
PC-DARTS+c/o [209]	CVPR20	3.6	97.43	0.1	1 1080Ti	
MiLeNAS[150]	CVPR20	3.87	97.66	0.3	-	
SGAS[210]	CVPR20	3.8	97.61	0.25	1 1080Ti	
GDAS-NSAS[211]	CVPR20	3.54	97.27	0.4	-	
NASBOT[157]	NeurIPS18	-	91.31	1.7	-	SMBO
PNAS [18]	ECCV18	3.2	96.59	225	-	
EPNAS[163]	BMVC18	6.6	96.29	1.8	1	
GHN[212]	ICLR19	5.7	97.16	0.84	-	RS
NAO+random+c/o[166]	NeurIPS18	10.6	97.52	200	200 V100	
SMASH [14]	ICLR18	16	95.97	1.5	-	
Hierarchical-random [19]	ICLR18	15.7	96.09	8	200	
RandomNAS [177]	UAI19	4.3	97.15	2.7	-	
DARTS - random+c/o [17]	ICLR19	3.2	96.71	4	1	
RandomNAS-NSAS[211]	CVPR20	3.08	97.36	0.7	-	GD+SMBO EA+RL EA+GD
NAO+weight sharing+c/o [166]	NeurIPS18	2.5	97.07	0.3	1 V100	
RENASNet+c/o[42]	CVPR19	3.5	91.12	1.5	4	
CARS[40]	CVPR20	3.6	97.38	0.4	-	