

Dive into Deep Learning

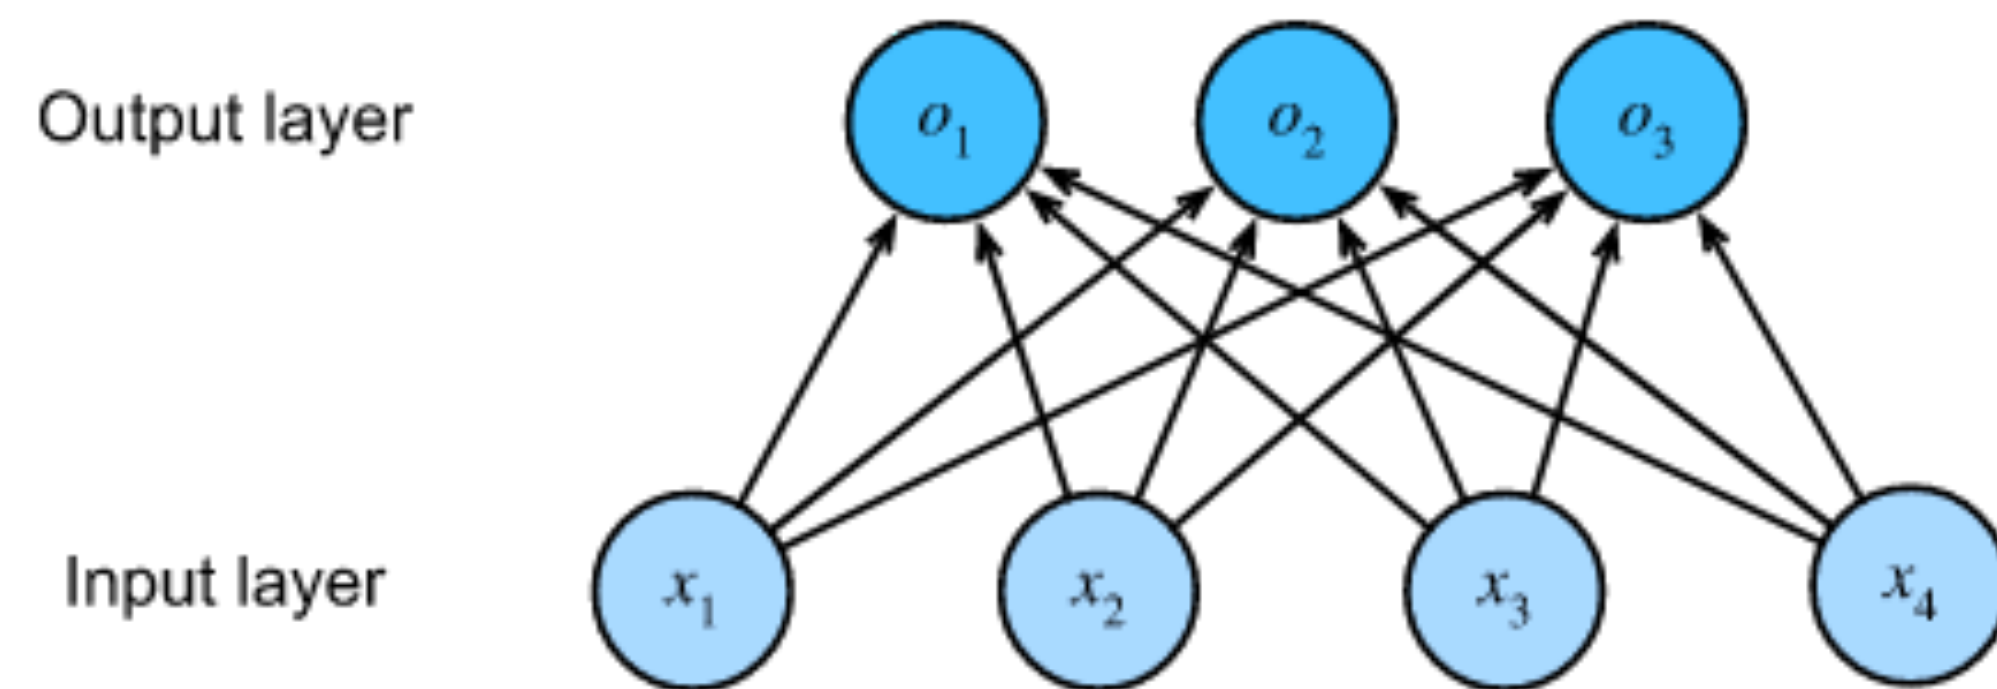
Chapter 4. Multilayer Perceptrons

Wayne L, Oct 18, 2020

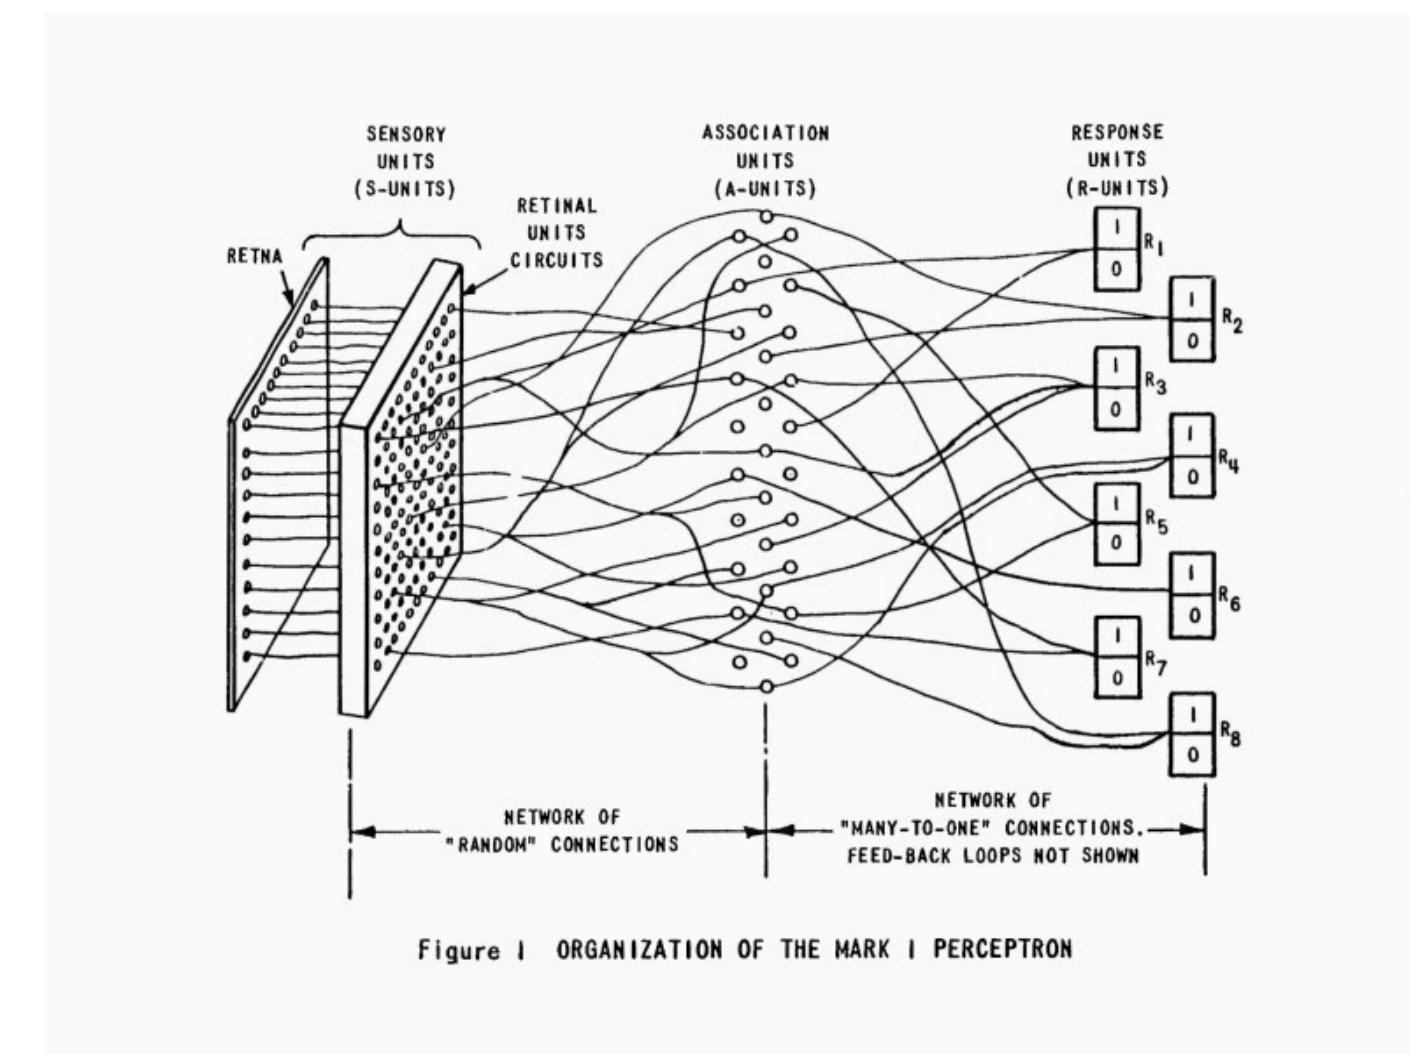
Previously

Perceptron (Linear = Single = Simple = Feedforward)

- This model mapped our inputs directly to our outputs via single affine transformation, followed by a softmax operation.



$$y = \mathbf{w}\mathbf{x} + b$$




Mark 1 Perceptron (1959), Frank Rosenblatt

Previously

AI History



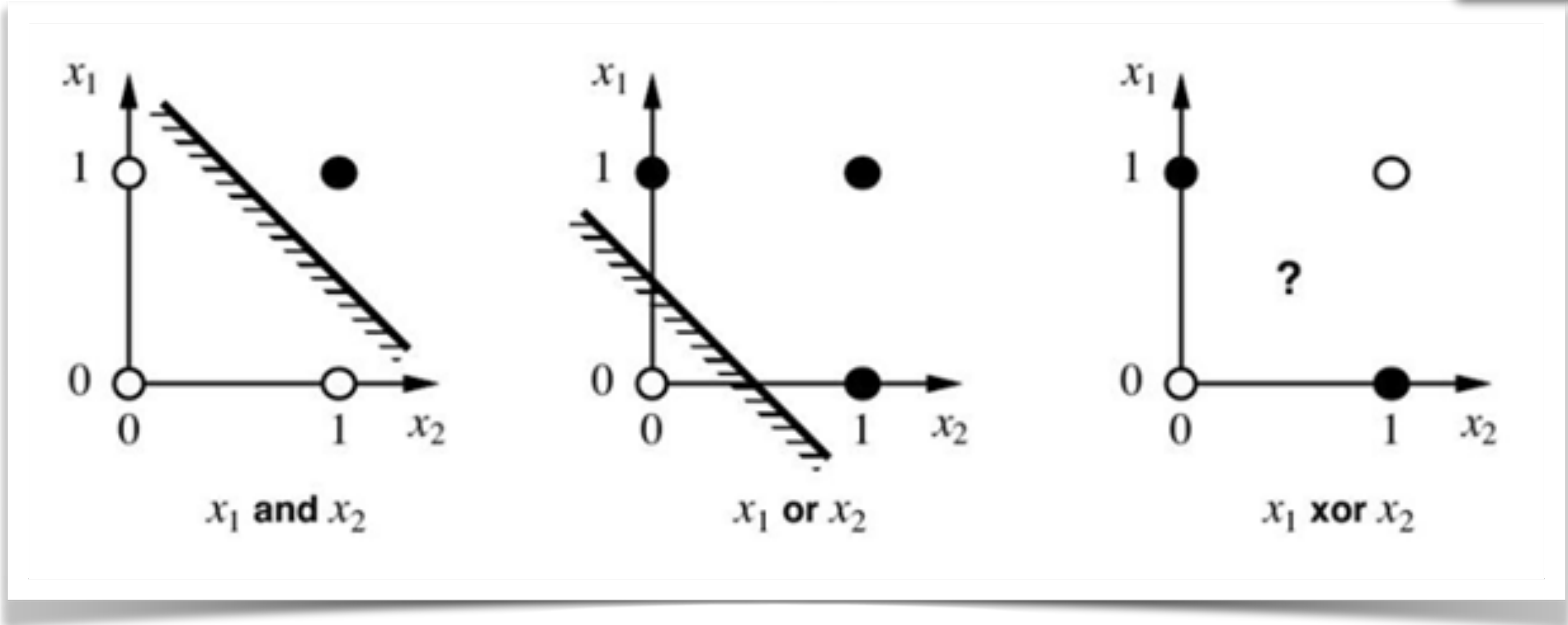

AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING” ...



| Timeline of AI Development |
|---|
| ▪ 1950s-1960s: First AI boom - the age of reasoning, prototype AI developed |
| ▪ 1970s: AI winter I |
| ▪ 1980s-1990s: Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making) |
| ▪ 1990s: AI winter II |
| ▪ 1997: Deep Blue beats Gary Kasparov |
| ▪ 2006: University of Toronto develops Deep Learning |
| ▪ 2011: IBM's Watson won Jeopardy |
| ▪ 2016: Go software based on Deep Learning beats world's champions |

XOR

Perceptrons (1969), Marvin Minsky




No one on earth had found a viable way to train

Multilayer Perceptrons

Perceptrons (Multi = Deep feedforward)

- The goal of a feedforward network is to approximate some function f^* .
- For example, for a classifier, $y = f^*(x)$ maps an input x to a category y .
A feedforward network defines a mapping $\mathbf{y} = f(\mathbf{x}; \theta)$ and learns the value of the parameters θ that result in the best function approximation.
- **Universal approximation theorem** means that regardless of what function we are trying to learn, we know that a large MLP will be able to represent this function.

$$f(x) = f^*(x) \approx y$$



$$\mathbf{O} = \mathbf{H}\mathbf{W}^{(2)} + \mathbf{b}^{(2)}$$

$$\mathbf{H} = \mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)}$$

$$\mathbf{W} = \mathbf{W}^{(1)}\mathbf{W}^{(2)}$$
$$\mathbf{b} = \mathbf{b}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)}$$

$$\mathbf{O} = (\mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)})\mathbf{W}^{(2)} + \mathbf{b}^{(2)} = \mathbf{X}\mathbf{W}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)} = \mathbf{X}\mathbf{W} + \mathbf{b}$$

Even though MLP is going deeper, it can be equivalent with single-layer model !

Multilayer Perceptrons

From Linear to Nonlinear

- In order to realize the potential of multilayer architectures, we need one more key ingredient: a nonlinear *activation function* σ to be more **expressive**.
- MLPs are **universal approximators**, however, it does not mean that we can solve all of problems with MLPs. In fact, we can approximate many functions much more compactly by using deeper (or wider) networks.
- Each neuron acts as a **linear SVM**, however, ...
 - its output is not interpreted immediately,
 - but it becomes a new feature,
 - to be forwarded to the next layer for further analysis **#SVMcascade**

$$\mathbf{H}^{(1)} = \sigma_1(\mathbf{X}\mathbf{W}^{(1)} + \mathbf{b}^{(1)})$$

$$\mathbf{H}^{(2)} = \sigma_2(\mathbf{H}^{(1)}\mathbf{W}^{(2)} + \mathbf{b}^{(2)})$$

Multilayer Perceptrons

Activation function

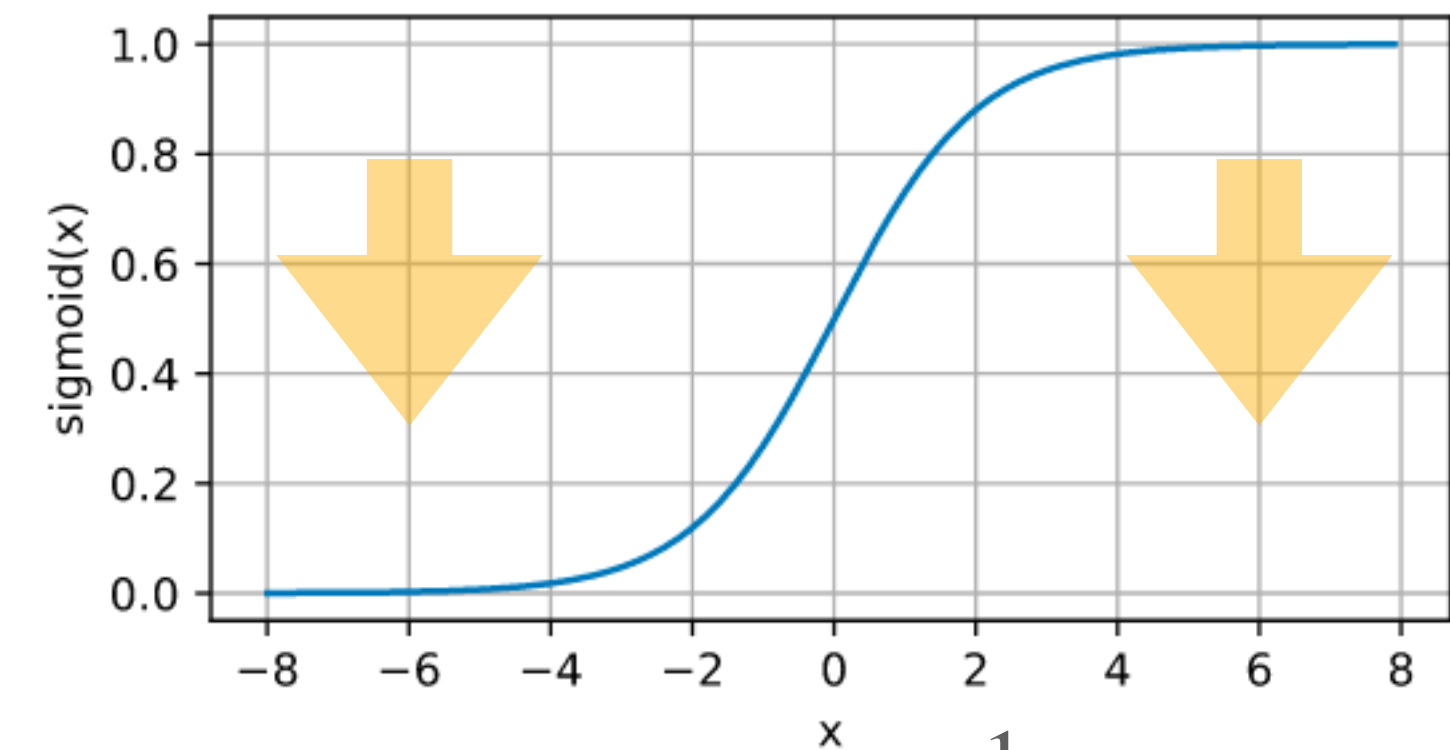
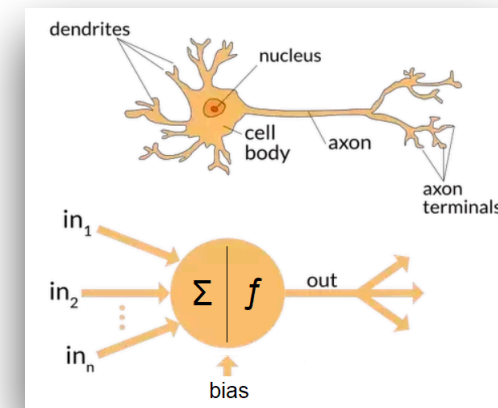
- **Activation function** decide whether a neuron should be activated or not by calculating the weighted sum and further adding bias with it.
- They are **differentiable** operators to transform input signal to outputs, while most of them add non-linearity.

| | |
|------------------------------------|--|
| <code>nn.ELU</code> | Applies the element-wise function: |
| <code>nn.Hardshrink</code> | Applies the hard shrinkage function element-wise: |
| <code>nn.Hardsigmoid</code> | Applies the element-wise function: |
| <code>nn.Hardtanh</code> | Applies the HardTanh function element-wise |
| <code>nn.Hardswish</code> | Applies the hardswish function, element-wise, as described in the paper: |
| <code>nn.LeakyReLU</code> | Applies the element-wise function: |
| <code>nn.LogSigmoid</code> | Applies the element-wise function: |
| <code>nn.MultiheadAttention</code> | Allows the model to jointly attend to information from different representation subspaces. |
| <code>nn.PReLU</code> | Applies the element-wise function: |
| <code>nn.ReLU</code> | Applies the rectified linear unit function element-wise: |
| <code>nn.ReLU6</code> | Applies the element-wise function: |
| <code>nn.RReLU</code> | Applies the randomized leaky rectified liner unit function, element-wise, as described in the paper: |
| <code>nn.SELU</code> | Applied element-wise, as: |
| <code>nn.CELU</code> | Applies the element-wise function: |
| <code>nn.GELU</code> | Applies the Gaussian Error Linear Units function: |
| <code>nn.Sigmoid</code> | Applies the element-wise function: |
| <code>nn.Softplus</code> | Applies the element-wise function: |
| <code>nn.Softshrink</code> | Applies the soft shrinkage function elementwise: |
| <code>nn.Softsign</code> | Applies the element-wise function: |
| <code>nn.Tanh</code> | Applies the element-wise function: |
| <code>nn.Tanhshrink</code> | Applies the element-wise function: |
| <code>nn.Threshold</code> | Thresholds each element of the input Tensor. |

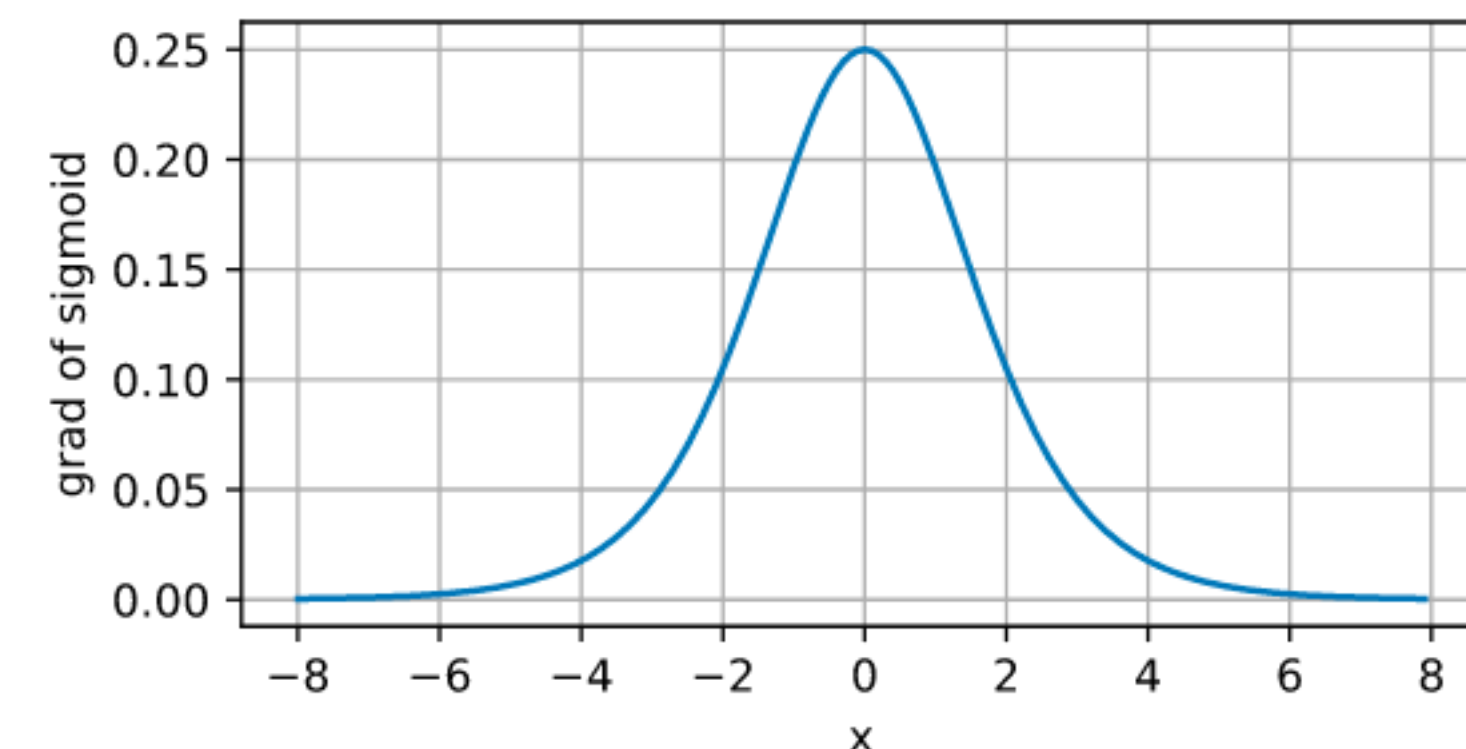
Multilayer Perceptrons

Sigmoid

- In the earliest neural networks, scientists were interested in modeling biological neurons which either fire or do not fire.
- When attention shifted to gradient based learning, the sigmoid function was a natural choice because it is a smooth, differentiable approximation to a thresholding unit.
- widely used as activation functions on the output units, when we want to interpret the outputs as probabilities for binary classification problems (special case of the softmax).



$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}$$




vanishing
gradient

$$\frac{d}{dx} \text{sigmoid}(x) = \frac{\exp(-x)}{(1 + \exp(-x))^2} = \text{sigmoid}(x)(1 - \text{sigmoid}(x))$$


Multilayer Perceptrons

Tanh (hyperbolic tangent)

- Like the sigmoid function, the tanh function also squashes its inputs, transforming them into elements on the interval between -1 and 1.
- Although the shape of the function is similar to that of the sigmoid function, the tanh function exhibits point symmetry about the origin of the coordinate system.



$$\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$$



$$\frac{d}{dx} \tanh(x) = 1 - \tanh^2(x)$$

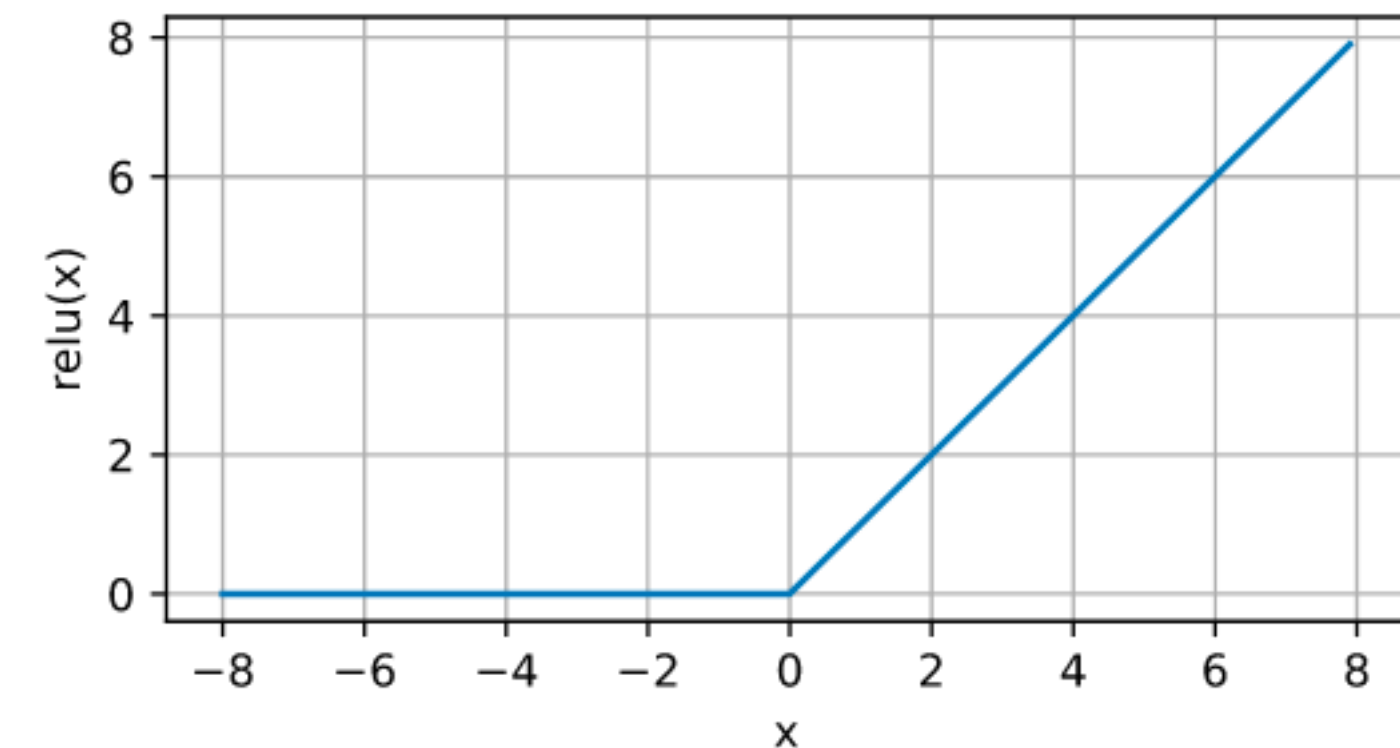
vanishing
gradient

Multilayer Perceptrons

ReLU (rectified linear unit)

- The reason for using ReLU is that its derivatives are particularly well behaved: either they vanish or they just let the argument through.
- This makes optimization better behaved and it mitigated the well-documented problem of vanishing gradients that plagued previous versions of neural networks.

$$pReLU(x) = \max(0, x) + \alpha \min(0, x)$$



$$ReLU(x) = \max(x, 0)$$

