

# Dive into DL

6.1 ~ 6.3

# Index

- 6. Convolutional Neural Networks
  - 6.1. From Fully-Connected Layers to Convolutions
  - 6.2. Convolutions for Images
  - 6.3. Padding and Stride

## 6.1. From Fully-Connected Layers to Convolutions

- fully-connected layer characterized: pixels x features
- ***translation invariance***: network should respond similarly to the same patch
- ***locality principle***: focus on local regions, without regard for the contents of the image in distant regions.



# 6.1. From Fully-Connected Layers to Convolutions

## 6.1.2. Constraining the MLP

$$\begin{aligned} [\mathbf{H}]_{i,j} &= [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l} \\ &= [\mathbf{U}]_{i,j} + \sum_a \sum_b [\mathbf{V}]_{i,j,a,b} [\mathbf{X}]_{i+a,j+b}, \end{aligned}$$

$\mathbf{X}$  : *Input*

$\mathbf{H}$  : hidden representations

$\mathbf{W}$  : fourth-order weight tensors

$\mathbf{U}$  : biases

re-index :  $k = i + a, l = j + b$   $[\mathbf{V}]_{i,j,a,b} = [\mathbf{W}]_{i,j,i+a,j+b}$ .

### 6.1.2.1. Translation Invariance

only possible if  $\mathbf{V}$  and  $\mathbf{U}$  do not actually depend on  $(i, j)$ , i.e., we have  $[\mathbf{V}]_{i,j,a,b} = [\mathbf{V}]_{a,b}$

$$[\mathbf{H}]_{i,j} = u + \sum_a \sum_b [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}.$$

### 6.1.2.2. Locality


$|a| > \Delta$  or  $|b| > \Delta$ , set  $[\mathbf{V}]_{a,b} = 0$ .

$$[\mathbf{H}]_{i,j} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}.$$


$\mathbf{V}$  is referred to as a *convolution kernel*, a *filter*, or simply the layer's *weights* that are often learnable parameters.

## Constraining the MLP

$$\begin{aligned}
 [\mathbf{H}]_{i,j} &= [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l} \\
 &= [\mathbf{U}]_{i,j} + \sum_a \sum_b [\mathbf{V}]_{i,j,a,b} [\mathbf{X}]_{i+a,j+b}.
 \end{aligned}$$



## Translation Invariance



# 6.1. From Fully-Connected Layers to Convolutions

## 6.1.2. Constraining the MLP

$$\begin{aligned} [\mathbf{H}]_{i,j} &= [\mathbf{U}]_{i,j} + \sum_k \sum_l [\mathbf{W}]_{i,j,k,l} [\mathbf{X}]_{k,l} \\ &= [\mathbf{U}]_{i,j} + \sum_a \sum_b [\mathbf{V}]_{i,j,a,b} [\mathbf{X}]_{i+a,j+b}, \end{aligned}$$

$\mathbf{X}$  : *Input*

$\mathbf{H}$  : hidden representations

$\mathbf{W}$  : fourth-order weight tensors

$\mathbf{U}$  : biases

re-index :  $k = i + a, l = j + b$   $[\mathbf{V}]_{i,j,a,b} = [\mathbf{W}]_{i,j,i+a,j+b}$ .

### 6.1.2.1. Translation Invariance

only possible if  $\mathbf{V}$  and  $\mathbf{U}$  do not actually depend on  $(i, j)$ , i.e., we have  $[\mathbf{V}]_{i,j,a,b} = [\mathbf{V}]_{a,b}$

$$[\mathbf{H}]_{i,j} = u + \sum_a \sum_b [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}.$$

### 6.1.2.2. Locality

$|a| > \Delta$  or  $|b| > \Delta$ , set  $[\mathbf{V}]_{a,b} = 0$ .

$$[\mathbf{H}]_{i,j} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}.$$

$\mathbf{V}$  is referred to as a *convolution kernel*, a *filter*, or simply the layer's *weights* that are often learnable parameters.

# 6.1. From Fully-Connected Layers to Convolutions

## 6.1.3. Convolutions

$$f, g: \mathbb{R}^d \rightarrow \mathbb{R}$$

$$(f * g)(\mathbf{x}) = \int f(\mathbf{z})g(\mathbf{x} - \mathbf{z})d\mathbf{z}.$$

$$(f * g)(i) = \sum_a f(a)g(i - a).$$

running over  $\mathbf{z}$



sum with indices  $(a, b)$  for  $f$  and  $(i - a, j - b)$  for  $g$ ,

$$(f * g)(i, j) = \sum_a \sum_b f(a, b)g(i - a, j - b).$$



## 6.1. From Fully-Connected Layers to Convolutions

- 6.1.4. “Where’s Waldo” Revisited

The convolutional layer picks windows of a given size and weighs intensities according to the filter  $V$



# 6.1. From Fully-Connected Layers to Convolutions

## 6.1.5. Channel

In reality, third-order tensors, characterized by a height, width, and channel  
3 channels: red, green, and blue

support multiple channels in both inputs  $X$  and hidden representations  $H$

$$[\mathbf{H}]_{i,j} = u + \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} [\mathbf{V}]_{a,b} [\mathbf{X}]_{i+a,j+b}.$$


$$[\mathbf{H}]_{i,j,d} = \sum_{a=-\Delta}^{\Delta} \sum_{b=-\Delta}^{\Delta} \sum_c [\mathbf{V}]_{a,b,c,d} [\mathbf{X}]_{i+a,j+b,c},$$

$d$  : indexes the output channels in the hidden representations

## 6.2. Convolutions for Images

### 6.2.1. The Cross-Correlation Operation

an input tensor and a kernel tensor are combined to produce an output tensor.



$$\begin{aligned}0 \times 0 + 1 \times 1 + 3 \times 2 + 4 \times 3 &= 19, \\1 \times 0 + 2 \times 1 + 4 \times 2 + 5 \times 3 &= 25, \\3 \times 0 + 4 \times 1 + 6 \times 2 + 7 \times 3 &= 37, \\4 \times 0 + 5 \times 1 + 7 \times 2 + 8 \times 3 &= 43.\end{aligned}$$

### 6.2.2. Convolutional Layers

A convolutional layer cross-correlates the input and kernel and adds a scalar bias to produce an output

```
from d2l import tensorflow as d2l
import tensorflow as tf

def corr2d(X, K): #@save
    """Compute 2D cross-correlation."""
    h, w = K.shape
    Y = tf.Variable(tf.zeros((X.shape[0] - h + 1, X.shape[1] - w + 1)))
    for i in range(Y.shape[0]):
        for j in range(Y.shape[1]):
            Y[i, j].assign(tf.reduce_sum(
                X[i:i + h, j:j + w] * K))
    return Y
```

```
X = tf.constant([[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
K = tf.constant([[0.0, 1.0], [2.0, 3.0]])
corr2d(X, K)
```

```
<tf.Variable 'Variable:0' shape=(2, 2) dtype=float32, numpy=
array([[19., 25.],
       [37., 43.]], dtype=float32)>
```

## 6.2. Convolutions for Images

### 6.2.3. Object Edge Detection in Images

construct an "image" of  $6 \times 8$  pixels

```
X = tf.Variable(tf.ones((6, 8)))  
X[:, 2:6].assign(tf.zeros(X[:, 2:6].shape))  
  
X  
  
<tf.Variable 'Variable:0' shape=(6, 8) dtype=float32, numpy=  
array([[1., 1., 0., 0., 0., 0., 1., 1.],  
       [1., 1., 0., 0., 0., 0., 1., 1.],  
       [1., 1., 0., 0., 0., 0., 1., 1.],  
       [1., 1., 0., 0., 0., 0., 1., 1.],  
       [1., 1., 0., 0., 0., 0., 1., 1.],  
       [1., 1., 0., 0., 0., 0., 1., 1.]])>
```

- construct a kernel K with a height of 1 and a width of 2.

```
K = tf.constant([[1.0, -1.0]])
```

$$Y = \text{corr2d}(X, K)$$

Y

```
<tf.Variable 'Variable_0' shape=(6, 7) dtype=float32, numpy=
array([[ 0.,  1.,  0.,  0.,  0., -1.,  0.],
       [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
       [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
       [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
       [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
       [ 0.,  1.,  0.,  0.,  0., -1.,  0.]])>
```

The kernel K only detects vertical edges.


`corr2d(tf.transpose(X), K)`[illegible]

## 6.3. Padding and Stride

### 6.3.1. Padding

when applying convolutional layers is that we tend to lose pixels on the perimeter of our image.

set the values of the extra pixels to zero.



$$0 \times 0 + 0 \times 1 + 0 \times 2 + 0 \times 3 = 0.$$

```
import tensorflow as tf
```

```
# We define a convenience function to calculate the convolutional layer. This  
# function initializes the convolutional layer weights and performs  
# corresponding dimensionality elevations and reductions on the input and  
# output
```

```
def comp_conv2d(conv2d, X):
```

```
    # Here (1, 1) indicates that the batch size and the number of channels  
    # are both 1
```

```
    X = tf.reshape(X, (1, ) + X.shape + (1, ))
```

```
    Y = conv2d(X)
```

```
    # Exclude the first two dimensions that do not interest us: examples and  
    # channels
```

```
    return tf.reshape(Y, Y.shape[1:3])
```

```
# Note that here 1 row or column is padded on either side, so a total of 2  
# rows or columns are added
```

```
conv2d = tf.keras.layers.Conv2D(1, kernel_size=3, padding='same')
```

```
X = tf.random.uniform(shape=(8, 8))
```

```
comp_conv2d(conv2d, X).shape
```

```
TensorShape([8, 8])
```

```
# Here, we use a convolution kernel with a height of 5 and a width of 3. The  
# padding numbers on either side of the height and width are 2 and 1,  
# respectively
```

```
conv2d = tf.keras.layers.Conv2D(1, kernel_size=(5, 3), padding='valid')
```

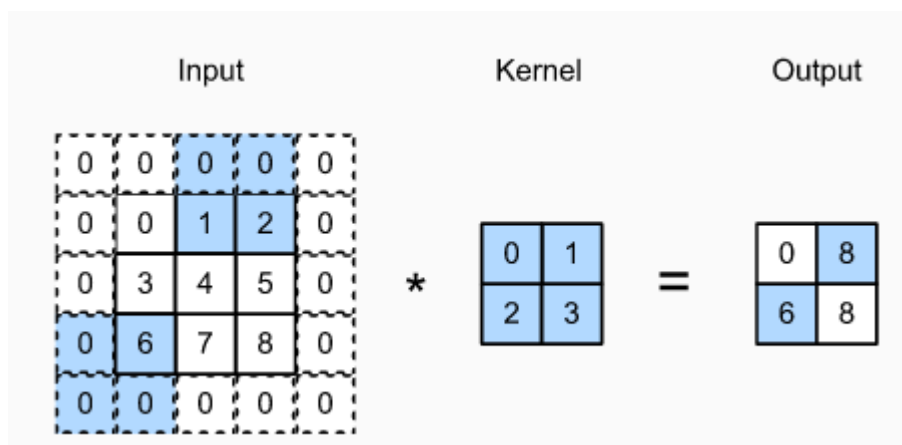
```
comp_conv2d(conv2d, X).shape
```

```
TensorShape([4, 6])
```

## 6.3. Padding and Stride

### 6.3.2. Stride

controls how depth columns around the spatial dimensions (width and height) are allocated



stride of 3 vertically and 2 horizontally

$$0 \times 0 + 0 \times 1 + 1 \times 2 + 2 \times 3 = 8, 0 \times 0 + 6 \times 1 + 0 \times 2 + 0 \times 3 = 6.$$

```
conv2d = tf.keras.layers.Conv2D(1, kernel_size=3, padding='same', strides=2)
comp_conv2d(conv2d, X).shape
```

```
TensorShape([4, 4])
```

```
conv2d = tf.keras.layers.Conv2D(1, kernel_size=3, padding='same', strides=2)
comp_conv2d(conv2d, X).shape
```

```
TensorShape([4, 4])
```