

Dive into Deep Learning


Chapter 8. Recurrent Neural Networks

Wayne L, Nov 15, 2020

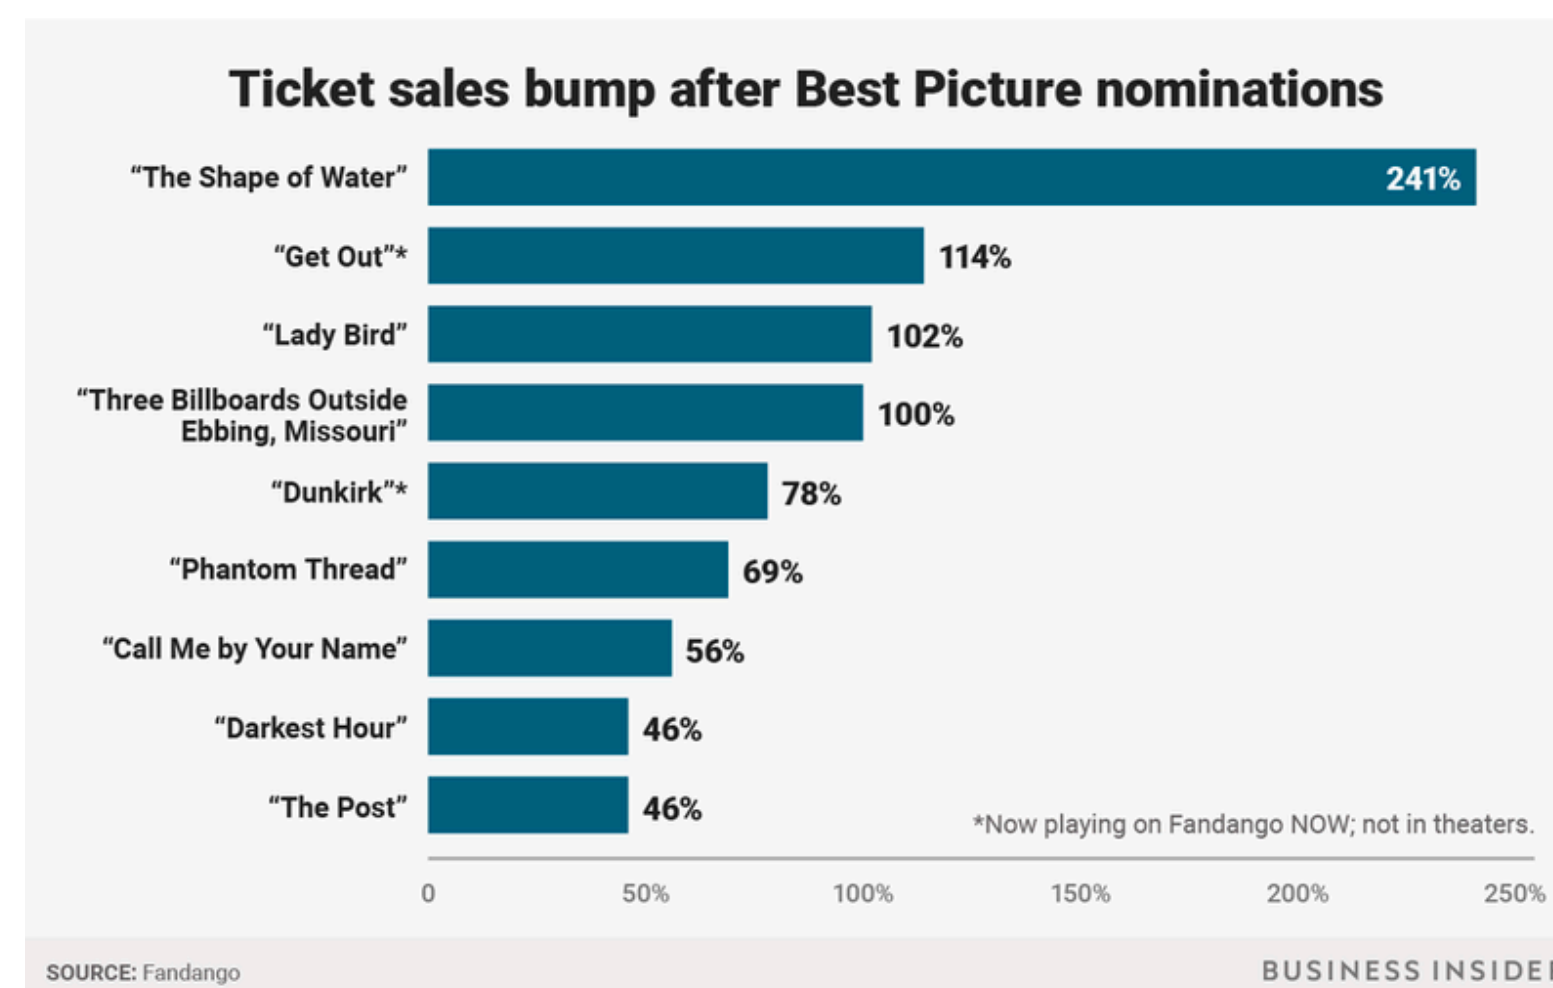
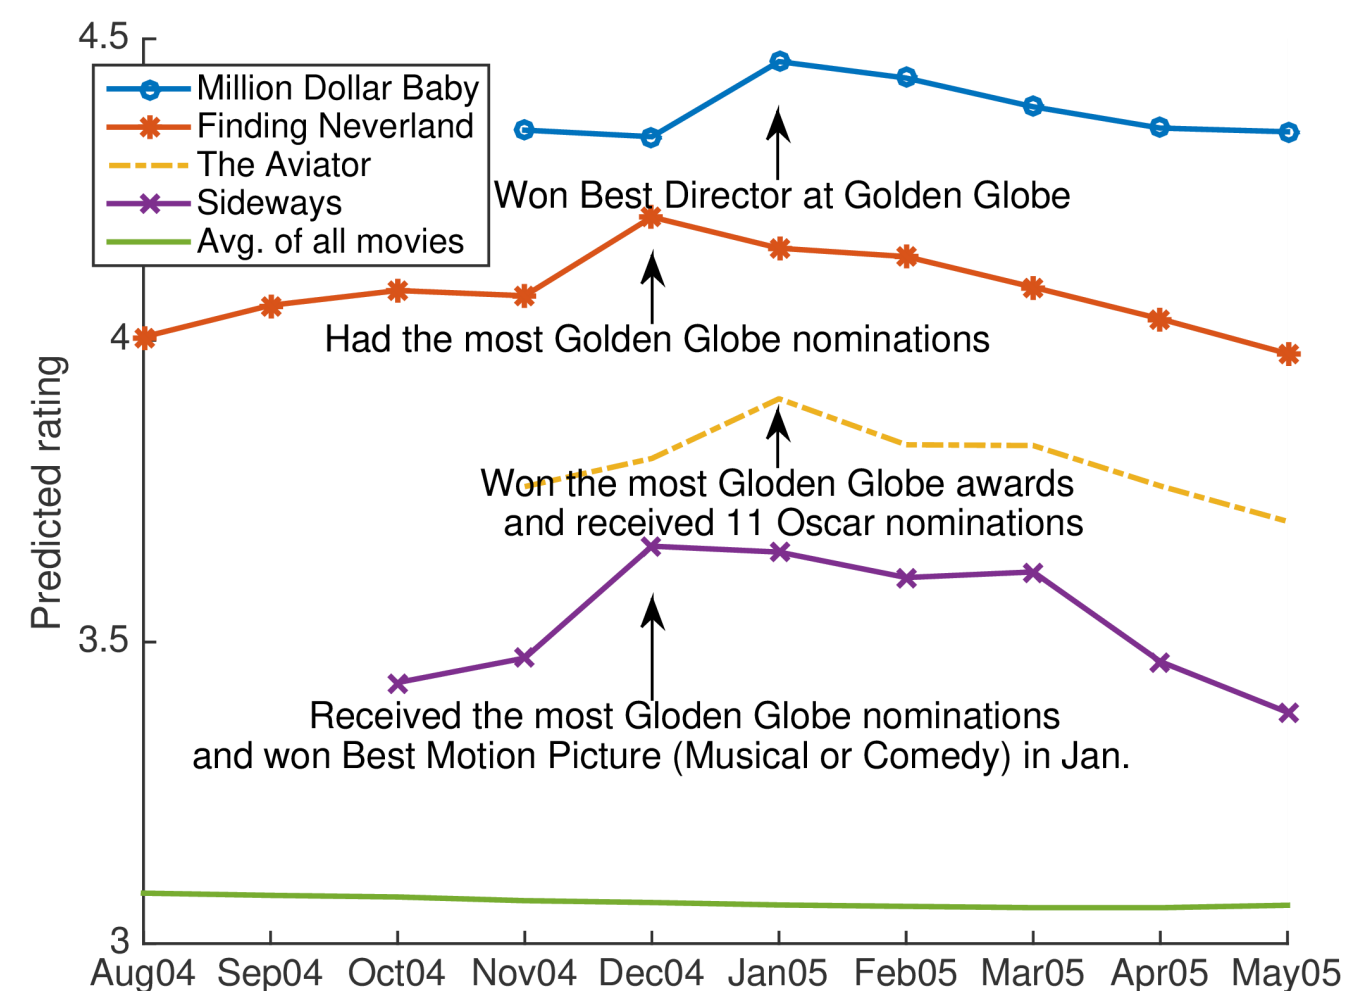
Sequence models

Oscar bump

- After the Oscar awards, ratings for the corresponding movie go up, even though it is still the same movie.
- This effect persists for a few months until the award is forgotten. It has been shown that the effect lifts rating by over half a point.



Wu, C.-Y., Ahmed, A., Beutel, A., Smola, A. J., & Jing, H. (2017). Recurrent recommender networks. *Proceedings of the tenth ACM international conference on web search and data mining* (pp. 495–503).



SBS CNBC PICK | 2020.02.12. | 네이버뉴스

오스카 후광에 '기생충 마케팅' 시작...글로벌 수입 주목


또 아카데미 수상으로 발생하는 단기 매출 급등 효과를 말하는 '오스카 뱀프'라는 용어가 있는데 요. 기생충에도 적용될 것으로 보입니다. 지난해 오스카 작품상을 수상한 영화 '그린북'이 수...



한국일보 PICK | 19면 1단 | 2020.02.10. | 네이버뉴스

'기생충'의 아카데미 수상 경제효과는? "오스카 뱀프 등 무궁무진"


기생충은 북미에서도 개봉한 지 4개월이 넘는 상태라 이미 영화를 본 관객이 많다면 오스카 뱀프 규모가 크지 않을 수 있다. 하지만 지금은 영화의 수익이 영화관에서 끝나지 않는 시대다. ...



연합뉴스TV | 2020.02.12. | 네이버뉴스

오스카 특수 '기생충'...온오프라인 휩쓴다

[리포터] 봉준호 감독의 영화 기생충이 오스카 4관왕을 차지한 뒤, 오스카 뱀프로 불리는 후광 효과를 누리고 있습니다. 워싱턴포스트가 미국 박스오피스 실적을 언급하며 아직 보지 못한 ...



YTN | 2020.02.19. | 네이버뉴스

[더뉴스-더쉬운경제] 아카데미 4관왕 '기생충'...경제적 효과는?

그것이 해당 영화뿐만 아니라 그해 만약에 외국일 경우에는 그 나라 전체의 문화산업 나아가서 국가 브랜드까지도 영향을 미친다고 해서 오스카 바운스, 오스카 뱀프 이야기가 많이 나오는...



매일경제 PICK | A2면 TOP | 2020.02.11. | 네이버뉴스

전세계 퍼지는 기생충 '품의 5억달러' 넘길까




Sequence models


Data usually didn't IID

- Various data we have


- Speech recognition
- Sentiment classification
- DNA sequence analysis
- Machine translation
- Action recognition
- NER
- ...



DNA sequence




Temperature sequence



Audio sequence

- Suppose that a trader who want to do well in the stock market on day t predicts x_t via $x_t \sim P(x_t \mid x_{t-1}, \dots, x_1)$.
(t : time, x_t : price at t)




Stock price

Autoregressive Models

- Autoregressive model
 - inputs: x_{t-1}, \dots, x_1
 - how to estimate $P(x_t | x_{t-1}, \dots, x_1)$ efficiently
- First, long sequence is not really necessary, instead, only use $x_{t-1}, \dots, x_{t-\tau}$ ($t > \tau$)
 - # of args always the same, allowing us to train a deep network
 - Such models will be called **autoregressive model**
- Secondly, to keep some summary h_t of the past observations, and at the same time update h_t in addition to the prediction $\hat{x}_t (=P(x_t | h_t))$
 - These models are also called **latent autoregressive models**

$$p(x) = p(x_1) \cdot p(x_2 | x_1) \cdot p(x_3 | x_1, x_2) \cdot \dots \cdot p(x_T | x_{T-\tau}, \dots, x_{T-1})$$




- In practice solve regression problem

$$\hat{x}_t = f(x_{t-\tau}, \dots, x_{t-1})$$

e.g. train an MLP on previously seen data

- Latent state summarizes all the relevant information about the past. So we get $h_t = f(x_1, \dots, x_{t-1}) = f(h_{t-1}, x_{t-1})$

$$p(h_t | h_{t-1}, x_{t-1}) \text{ and } p(x_t | h_t, x_{t-1})$$



Markov models

Next observation only depends on the past few terms

- First-order Markov model (if $\tau = 1$)

$$P(x_1, \dots, x_T) = \prod_{t=1}^T P(x_t | x_{t-1}) \text{ where } P(x_1 | x_0) = P(x_1)$$

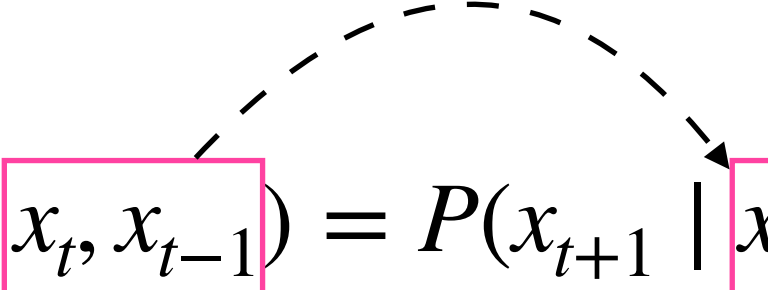
- Such models are particularly nice whenever x_t assumes **only a discrete value**, since in this case dynamic programming can be used to compute values along the chain exactly.

For instance, we can compute $P(x_{t+1} | x_{t-1})$ efficiently:

$$\begin{aligned} P(x_{t+1} | x_{t-1}) &= \frac{\sum_{x_t} P(x_{t+1}, x_t, x_{t-1})}{P(x_{t-1})} \\ &= \frac{\sum_{x_t} P(x_{t+1} | x_t, x_{t-1}) P(x_t, x_{t-1})}{P(x_{t-1})} & \frac{P(x_t \cap x_{t-1})}{P(x_{t-1})} = P(x_t | x_{t-1}) \\ &= \sum_{x_t} P(x_{t+1} | x_t) P(x_t | x_{t-1}) \end{aligned}$$

- First-order Markov model (if $\tau = 1$)

by using the fact that we only need to take into account a very short history of past observations


$$P(x_{t+1} | x_t, x_{t-1}) = P(x_{t+1} | x_t)$$


Casualty

It's clear that future events cannot influence the past


- Causality (physics) prevents the reverse direction
- 'wrong' direction often much more complex to model
- For instance, it has been shown that in some cases we can find $x_{t+1} = f(x_t) + \epsilon$ for some additive noise ϵ , whereas the converse is not true
- Typically forward direction that we're interested in estimating !

Hoyer, P. O., Janzing, D., Mooij, J. M., Peters, J., & Schölkopf, B. (2009). Nonlinear causal discovery with additive noise models. *Advances in neural information processing systems* (pp. 689–696).

$$p(x) = p(x_1) \cdot p(x_2 | x_1) \cdot p(x_3 | x_1, x_2) \cdot \dots \cdot p(x_T | x_1, \dots, x_{T-1})$$




$$p(x) = p(x_T) \cdot p(x_{T-1} | x_T) \cdot p(x_{T-2} | x_{T-1}, x_T) \cdot \dots \cdot p(x_1 | x_2, \dots, x_T)$$



Training sequence model

```
[3] T = 1000 # Generate a total of 1000 points
time = torch.arange(1, T + 1, dtype=torch.float32)
x = torch.sin(0.01 * time) + torch.normal(0, 0.2, (T,))
d2l.plot(time, [x], 'time', 'x', xlim=[1, 1000], figsize=(6, 3))
```



```
[4] tau = 4
```

embedding dimension = 4

```
features = torch.zeros((T - tau, tau))
for i in range(tau):
    features[:, i] = x[i: T - tau + i]
labels = d2l.reshape(x[tau:], (-1, 1))
```

```
[5] batch_size, n_train = 16, 600
# Only the first `n_train` examples are used for training
train_iter = d2l.load_array((features[:n_train], labels[:n_train]),
                           batch_size, is_train=True)
```

```
[6] # Function for initializing the weights of the network
def init_weights(m):
    if type(m) == nn.Linear:
        torch.nn.init.xavier_uniform_(m.weight)

# A simple MLP
def get_net():
    net = nn.Sequential(nn.Linear(4, 10),
                        nn.ReLU(),
                        nn.Linear(10, 1))
    net.apply(init_weights)
    return net

# Square loss
loss = nn.MSELoss()
```

MLP with 2 FC layers, ReLU and squared loss

```
[7] def train(net, train_iter, loss, epochs, lr):
    trainer = torch.optim.Adam(net.parameters(), lr)
    for epoch in range(epochs):
        for X, y in train_iter:
            trainer.zero_grad()
            l = loss(net(X), y)
            l.backward()
            trainer.step()
        print(f'epoch {epoch + 1}, '
              f'loss: {d2l.evaluate_loss(net, train_iter, loss):f}')


net = get_net()
train(net, train_iter, loss, 5, 0.01)
```

```
epoch 1, loss: 0.054968
epoch 2, loss: 0.055680
epoch 3, loss: 0.058150
epoch 4, loss: 0.049353
epoch 5, loss: 0.050014
```


Training sequence model

1-step-ahead prediction

```
[8] onestep_preds = net(features)
    d2l.plot([time, time[tau:]], [d2l.numpy(x), d2l.numpy(onestep_preds)], 'time',
            'x', legend=['data', '1-step preds'], xlim=[1, 1000],
            figsize=(6, 3))
```




k-step-ahead prediction

```
[11] max_steps = 64
```

```
[12] features = torch.zeros((T - tau - max_steps + 1, tau + max_steps))
    # Column `i` (`i` < `tau`) are observations from `x` for time steps from
    # `i + 1` to `i + T - tau - max_steps + 1`
    for i in range(tau):
        features[:, i] = x[i: i + T - tau - max_steps + 1].T

    # Column `i` (`i` >= `tau`) are the (`i - tau + 1`)-step-ahead predictions for
    # time steps from `i + 1` to `i + T - tau - max_steps + 1`
    for i in range(tau, tau + max_steps):
        features[:, i] = d2l.reshape(net(features[:, i - tau: i]), -1)
```

```
[13] steps = (1, 4, 16, 64)
    d2l.plot([time[tau + i - 1: T - max_steps + i] for i in steps],
            [d2l.numpy(features[:, tau + i - 1]) for i in steps], 'time', 'x',
            legend=[f'{i}-step preds' for i in steps], xlim=[5, 1000],
            figsize=(6, 3))
```



Text Processing

Basic Idea - map text into sequence of IDs

- **Character Encoding** (each character has one ID)
 - Small vocabulary
 - Doesn't work so well (DNN needs to learn spelling)
- **Word Encoding** (each word has one ID)
 - Accurate spelling
 - Doesn't work so well (huge vocabulary = costly multinomial)
- **Byte Pair Encoding** (Goldilocks zone)
 - Frequent subsequences (like syllables)

Language Models and the Dataset

- Given x_1, x_2, \dots, x_T (text sequence of length T), the goal of a language model is to estimate the joint probability of the sequence $P(x_1, x_2, \dots, x_T)$
- For instance, an ideal language model would be able to generate natural text just on its own, simply by drawing one token at a time $x_t \sim P(x_t \mid x_{t-1}, \dots, x_1)$.
- LMs are of great service even in their limited form.
- the phrases “to recognize speech” and “to wreck a nice beach” sound very similar. (ambiguity)
- “dog bites man” vs. “man bites dog”
“I want to eat grandma” vs. “I want to eat, grandma”

Language Models and the Dataset

- Tokenize text data at the word level, applying basic probability rules

$$p(w_1, w_2, \dots, w_T) = \prod_{t=1}^T p(w_t | w_1, \dots, w_{t-1})$$

- $p(\text{Statistics, is, fun, .})$
 $= p(\text{Statistics})p(\text{is} | \text{Statistics})p(\text{fun} | \text{Statistics, is})p(. | \text{Statistics, is, fun})$
- In order to compute the LM, we need to calculate probability of words and the conditional probability of a word given the previous few words.

- $\hat{p}(\text{is} | \text{Statistics}) = \frac{n(\text{Statistics, is})}{n(\text{Statistics})}$


($n(x)$ and $n(x, n')$ are the number of occurrences of singletons and consecutive word pairs)

- **Laplace smoothing**

$$\hat{p}(w) = \frac{n(w) + \epsilon_1/m}{n + \epsilon_1}$$

$$\hat{p}(w' | w) = \frac{n(w, w') + \epsilon_2 \hat{p}(w')}{n(w) + \epsilon_2}$$

$$\hat{p}(w'' | w', w) = \frac{n(w, w', w'') + \epsilon_3 \hat{p}(w', w'')}{n(w, w') + \epsilon_3}$$



Zipf's law: $n_i \propto \frac{1}{i^\alpha}$,